

LECTURE 17

Synchronization without Locks

Xuhao Chen

November 25, 2025



Outline

- Sequential Consistency (SC)
 - Mutual Exclusion w/o Locks under SC
- Relaxed Memory Consistency
 - Total Store Order (TSO)
- Compare and Swap (CAS)
- Lock-free Algorithms
 - Lock-free Stack
 - The ABA Problem

Memory Models

Initially, $a = b = 0$.

Processor 0

```
mov 1, a      ;Store  
mov b, %ebx   ;Load
```

Processor 1


```
mov 1, b      ;Store  
mov a, %eax   ;Load
```

- Q. Is it possible that Processor 0's $\%ebx$ and Processor 1's $\%eax$ both contain the value 0 after the processors have both executed their code?
- A. It depends on the **memory model**: how memory operations behave in the parallel computer system.

Sequential Consistency

“The result of any execution is the same as if the operations of all the processors were executed in some sequential order, and the operations of each individual processor appear in this sequence in the order specified by its program.”

— Leslie Lamport [1979]



- The sequence of instructions as defined by a processor’s program are **interleaved** with the corresponding sequences defined by the other processors’ programs to produce a global **linear order** of all instructions.

SC Example

Initially, $a = b = 0$.

Processor 0

1 `mov 1, a` ;Store
2 `mov b, %ebx` ;Load

Processor 1

3 `mov 1, b` ;Store
4 `mov a, %eax` ;Load


	Interleavings					
	1	1	1	3	3	3
	2	3	3	1	1	4
	3	2	4	2	4	1
	4	4	2	4	2	2
%eax	1	1	1	1	1	0
%ebx	0	1	1	1	1	1

Sequential consistency implies that no execution ends with $\%eax = \%ebx = 0$.

Sequential Consistency

“The result of any execution is the same as if the operations of all the processors were executed in some sequential order, and the operations of each individual processor appear in this sequence in the order specified by its program.”

— Leslie Lamport [1979]



- The sequence of instructions as defined by a processor’s program are **interleaved** with the corresponding sequences defined by the other processors’ programs to produce a global **linear order** of all instructions.
- A **LOAD** instruction receives the value stored to that address by the most recent **STORE** instruction that precedes the **LOAD**, according to the linear order.
- The hardware can do whatever it wants, but for the execution to be sequentially consistent, it must **appear** as if **LOAD**’s and **STORE**’s obey some global linear order.

Reasoning about Sequential Consistency

- An execution induces a “happens before” relation, denoted as $\rightarrow$.
- The $\rightarrow$ relation is linear (a total order), meaning that in particular for any two distinct instructions x and y , either $x \rightarrow y$ or $y \rightarrow x$.
- The $\rightarrow$ relation respects processor order, the order of instructions in each processor.
- A **LOAD** from a location in memory reads the value written by the most recent **STORE** to that location according to $\rightarrow$.
- For the memory resulting from an execution to be sequentially consistent, there must exist such a linear order $\rightarrow$ that yields that memory state.

Outline

- Sequential Consistency (SC)
 - ▣ **Mutual Exclusion w/o Locks under SC**
- Relaxed Memory Consistency
 - ▣ Total Store Order (TSO)
- Compare and Swap (CAS)
- Lock-free Algorithms
 - ▣ Lock-free Stack
 - ▣ The ABA Problem

Mutual-Exclusion Problem

Recall

- A **critical section** is a piece of code that accesses a shared data structure that must not be executed by two or more parallel strands (**mutual exclusion**)
- Computer hardware provides **atomic read-modify-write** instructions
 - e.g., atomic swap (X86 **xchg**), **TEST-AND-SET**, **COMPARE-AND-SWAP**, **LOAD-LINKED-STORE-CONDITIONAL**
- Synchronization libraries use these instructions to implement locks, but you can use them directly
 - The C library **stdatomic.h*** provides a long list of atomics for most architectures
 - LLVM & GCC provide compiler built-in functions for synchronization, but less **portable**


*See <http://en.cppreference.com/w/c/atomic>.

Mutual-Exclusion Problem

- Q. Can mutual exclusion be implemented with atomic **LOAD**'s and **STORE**'s as the only memory operations?
- A. Yes, **Theodorus J. Dekker** and **Edsger Dijkstra** showed that it can, at least for computers with sequentially consistent memory models.



Peterson's Algorithm



```
widget x; //protected variable  
bool A_wants = false;  
bool B_wants = false;  
enum {A, B} turn;
```

Alice

```
A_wants = true;  
turn = B;  
while (B_wants && turn==B);  
frob(&x); //critical section  
A_wants = false;
```

Bob

```
B_wants = true;  
turn = A;  
while (A_wants && turn==A);  
borf(&x); //critical section  
B_wants = false;
```

Peterson's Algorithm

Alice

Bob

```
1 A_wants = true;  
2 turn = B;  
  while (B_wants && turn==B);  
  frob(&x); //critical section  
  A_wants = false;
```

```
3 B_wants = true;  
4 turn = A;  
  while (A_wants && turn==A);  
  borf(&x); //critical section  
  B_wants = false;
```

Before Alice enters its `while` loop, if the interleaving was **1 → 2 → 3 → 4**

Then **B_wants** is **true**

turn is **A**

So Alice goes into the critical section.

In the meantime, before Bob enters its `while` loop, and before Alice exit,

A_wants is **true**

So Bob waits, until Alice exits and sets **A_wants** to be **false**.

Proof of Mutual Exclusion

Theorem. Peterson's algorithm guarantees mutual exclusion.

Proof.

- Assume by way of contradiction that both Alice and Bob execute the critical section concurrently.
- Let's derive a contradiction.

Proof of Mutual Exclusion

Alice

```
1 A_wants = true;  
2 turn = B;  
→ while (B_wants && turn==B);  
   frob(&x); //critical section  
   A_wants = false;
```

Bob

```
3 B_wants = true;  
4 turn = A;  
5 while (A_wants && turn==A);  
→ borf(&x); //critical section  
   B_wants = false;
```

- Consider **the most-recent time** that each of them executed the code before entering the critical section.

Proof of Mutual Exclusion

Alice

2

```
A_wants = true;  
turn = B;  
while (B_wants && turn==B);  
frob(&x); //critical section  
A_wants = false;
```

Bob

4

```
B_wants = true;  
turn = A;  
while (A_wants && turn==A);  
borf(&x); //critical section  
B_wants = false;
```

- Assume **WLOG** that **Bob** was the last to write to **turn**:

$\text{write}_A(\text{turn} = B) \rightarrow \text{write}_B(\text{turn} = A)$

$2 \rightarrow 4$

Proof of Mutual Exclusion

Alice

```
1 A_wants = true;  
2 turn = B;  
while (B_wants && turn==B);  
frob(&x); //critical section  
A_wants = false;
```

Bob

```
B_wants = true;  
turn = A;  
while (A_wants && turn==A);  
borf(&x); //critical section  
B_wants = false;
```

- Assume WLOG that Bob was the last to write to **turn**:

$\text{write}_A(\text{turn} = B) \rightarrow \text{write}_B(\text{turn} = A) \quad 2 \rightarrow 4$

- Alice's program order:

$\text{write}_A(A_wants = \text{true}) \rightarrow \text{write}_A(\text{turn} = B) \quad 1 \rightarrow 2$

Proof of Mutual Exclusion

Alice

```
A_wants = true;  
turn = B;  
while (B_wants && turn==B);  
frob(&x); //critical section  
A_wants = false;
```

Bob

```
B_wants = true;  
turn = A;  
while (A_wants && turn==A);  
borf(&x); //critical section  
B_wants = false;
```

4

5

- Assume WLOG that Bob was the last to write to **turn**:

$\text{write}_A(\text{turn} = B) \rightarrow \text{write}_B(\text{turn} = A) \quad 2 \rightarrow 4$

- Alice's program order:

$\text{write}_A(A_wants = \text{true}) \rightarrow \text{write}_A(\text{turn} = B) \quad 1 \rightarrow 2$

- Bob's program order:

$\text{write}_B(\text{turn} = A) \rightarrow \text{read}_B(A_wants) \rightarrow \text{read}_B(\text{turn}) \quad 4 \rightarrow 5$

Proof of Mutual Exclusion

Alice

```
1 A_wants = true;  
2 turn = B;  
while (B_wants && turn==B);  
frob(&x); //critical section  
A_wants = false;
```

Bob

```
3 B_wants = true;  
4 turn = A;  
5 while (A_wants && turn==A);  
borf(&x); //critical section  
B_wants = false;
```

$\text{write}_A(\text{turn} = B) \rightarrow \text{write}_B(\text{turn} = A) \quad 2 \rightarrow 4$

$\text{write}_A(A_wants = \text{true}) \rightarrow \text{write}_A(\text{turn} = B) \quad 1 \rightarrow 2$

$\text{write}_B(\text{turn} = A) \rightarrow \text{read}_B(A_wants) \rightarrow \text{read}_B(\text{turn}) \quad 4 \rightarrow 5$

Proof of Mutual Exclusion

Alice

```
1 A_wants = true;  
2 turn = B;  
while (B_wants && turn==B);  
frob(&x); //critical section  
A_wants = false;
```

Bob

```
3 B_wants = true;  
4 turn = A;  
5 while (A_wants && turn==A);  
borf(&x); //critical section  
B_wants = false;
```

$\text{write}_A(\text{A_wants} = \text{true}) \rightarrow \text{write}_A(\text{turn} = \text{B})$ 1 $\rightarrow$ 2

$\text{write}_A(\text{turn} = \text{B}) \rightarrow \text{write}_B(\text{turn} = \text{A})$ 2 $\rightarrow$ 4

$\text{write}_B(\text{turn} = \text{A}) \rightarrow \text{read}_B(\text{A_wants}) \rightarrow \text{read}_B(\text{turn})$ 4 $\rightarrow$ 5

Proof of Mutual Exclusion

Alice

```
1 A_wants = true;  
2 turn = B;  
while (B_wants && turn==B);  
frob(&x); //critical section  
A_wants = false;
```

Bob

```
3 B_wants = true;  
4 turn = A;  
5 while (A_wants && turn==A);  
borf(&x); //critical section  
B_wants = false;
```

$\text{write}_A(\text{A_wants} = \text{true}) \rightarrow \text{write}_A(\text{turn} = \text{B}) \quad 1 \rightarrow 2$

$\text{write}_A(\text{turn} = \text{B}) \rightarrow \text{write}_B(\text{turn} = \text{A}) \quad 2 \rightarrow 4$

$\text{write}_B(\text{turn} = \text{A}) \rightarrow \text{read}_B(\text{A_wants}) \rightarrow \text{read}_B(\text{turn}) \quad 4 \rightarrow 5$

Proof of Mutual Exclusion

Alice

```
1 A_wants = true;  
2 turn = B;  
while (B_wants && turn==B);  
frob(&x); //critical section  
A_wants = false;
```

Bob

```
3 B_wants = true;  
4 turn = A;  
5 while (A_wants && turn==A);  
borf(&x); //critical section  
B_wants = false;
```

$\text{write}_A(\text{A_wants} = \text{true}) \rightarrow$

$1 \rightarrow 2$

$\text{write}_A(\text{turn} = \text{B}) \rightarrow \text{write}_B(\text{turn} = \text{A})$

$2 \rightarrow 4$

$\text{write}_B(\text{turn} = \text{A}) \rightarrow \text{read}_B(\text{A_wants}) \rightarrow \text{read}_B(\text{turn})$

$4 \rightarrow 5$

Proof of Mutual Exclusion

Alice

```
1 A_wants = true;  
2 turn = B;  
while (B_wants && turn==B);  
frob(&x); //critical section  
A_wants = false;
```

Bob

```
3 B_wants = true;  
4 turn = A;  
5 while (A_wants && turn==A);  
borf(&x); //critical section  
B_wants = false;
```

$\text{write}_A(\text{A_wants} = \text{true}) \rightarrow$

$1 \rightarrow 2$

$\text{write}_A(\text{turn} = \text{B}) \rightarrow \text{write}_B(\text{turn} = \text{A})$

$2 \rightarrow 4$

$\text{write}_B(\text{turn} = \text{A}) \rightarrow \text{read}_B(\text{A_wants}) \rightarrow \text{read}_B(\text{turn})$

$4 \rightarrow 5$

Proof of Mutual Exclusion

Alice

```
1 A_wants = true;  
2 turn = B;  
while (B_wants && turn==B);  
frob(&x); //critical section  
A_wants = false;
```

Bob

```
3 B_wants = true;  
4 turn = A;  
5 while (A_wants && turn==A);  
borf(&x); //critical section  
B_wants = false;
```

$\text{write}_A(\text{A_wants} = \text{true}) \rightarrow$

$1 \rightarrow 2$

$\text{write}_A(\text{turn} = \text{B}) \rightarrow$

$2 \rightarrow 4$

$\text{write}_B(\text{turn} = \text{A}) \rightarrow \text{read}_B(\text{A_wants}) \rightarrow \text{read}_B(\text{turn})$

$4 \rightarrow 5$

$1 \rightarrow 2 \rightarrow 4 \rightarrow 5$

Proof of Mutual Exclusion

Alice

```
1 A_wants = true;  
2 turn = B;  
while (B_wants && turn==B);  
frob(&x); //critical section  
A_wants = false;
```

Bob

```
3 B_wants = true;  
4 turn = A;  
5 while (A_wants && turn==A);  
borf(&x); //critical section  
B_wants = false;
```

$\text{write}_A(\text{A_wants} = \text{true}) \rightarrow$

$\text{write}_A(\text{turn} = \text{B}) \rightarrow$

$\text{write}_B(\text{turn} = \text{A}) \rightarrow \text{read}_B(\text{A_wants}) \rightarrow \text{read}_B(\text{turn})$

- What did Bob read before entering critical section given $1 \rightarrow 2 \rightarrow 4 \rightarrow 5$

$\left. \begin{array}{l} \text{A_wants: true} \\ \text{turn: A} \end{array} \right\} \text{Bob should spin. Contradiction. } \blacksquare$

Deadlock Freedom

Alice

```
A_wants = true;  
turn = B;  
while (B_wants && turn==B);  
frob(&x); //critical section  
A_wants = false;
```

Bob

```
B_wants = true;  
turn = A;  
while (A_wants && turn==A);  
borf(&x); //critical section  
B_wants = false;
```

Theorem: Peterson's algorithm guarantees **Deadlock Freedom**.

- If **Alice** and **Bob** both try to enter the critical section, then whoever writes last to **turn** spins and the other progresses.
- If only **Alice** tries to enter the critical section, then she progresses, since **B_wants** is false.
- If only **Bob** tries to enter the critical section, then he progresses, since **A_wants** is false.


Outline

- Sequential Consistency (SC)
 - Mutual Exclusion w/o Locks under SC
- **Relaxed Memory Consistency**
 - Total Store Order (TSO)
- Compare and Swap (CAS)
- Lock-free Algorithms
 - Lock-free Stack
 - The ABA Problem

Memory Models Today


- No modern-day processor implements sequential consistency
- All implement some form of **relaxed consistency**.
- Hardware actively reorders instructions.
- Compilers may reorder instructions too.

Instruction Reordering



- Q. Why might the hardware or compiler decide to reorder these instructions?
- A. To obtain higher performance by covering load latency — instruction-level parallelism.

Instruction Reordering




Q. When is it safe for the hardware or compiler to perform this reordering?

A. When $a \neq b$.

A'. And there's no concurrency.

Hardware Reordering



- The processor can issue **STORE**'s faster than the network can handle them $\Rightarrow$ **store buffer**.
- Since a **LOAD** can stall the processor until it is satisfied, **loads take priority**, bypassing the store buffer.
- If a **LOAD** address matches an address in the store buffer, the store buffer returns the result.
- Thus, a **LOAD** can **bypass** a **STORE** to a different address.

x86-64 Total Store Order (TSO)

Locally:

Instruction Trace



1. **LOAD**'s are **not** reordered with **LOAD**'s.
2. **STORE**'s are **not** reordered with **STORE**'s.
3. **STORE**'s are **not** reordered with **prior LOAD**'s.

x86-64 Total Store Order (TSO)

Locally:

Instruction Trace

Store1

Store2

Load1

Load2

Store3

Store4

Load3

Load4

Load5

L
O
A
D
S

1. **LOAD**'s are **not** reordered with **LOAD**'s.
2. **STORE**'s are **not** reordered with **STORE**'s.
3. **STORE**'s are **not** reordered with **prior LOAD**'s.
4. A **LOAD** may be reordered with a prior **STORE** to a **different** location but **not** with a prior **STORE** to the **same** location.
5. **LOAD**'s and **STORE**'s are **not** reordered with **LOCK** instructions.

x86-64 Total Store Order (TSO)

Locally:

Instruction Trace

Store1
Store2
Load1
Load2
Store3
Store4
Load3
Load4
Load5

1. **LOAD**'s are **not** reordered with **LOAD**'s.
 2. **STORE**'s are **not** reordered with **STORE**'s.
 3. **STORE**'s are **not** reordered with **prior LOAD**'s.
 4. A **LOAD** is not reordered with a **STORE** to a **different** location.
 5. **LOCK** instructions respect a **global total order**.
- Total Store Ordering (TSO) is weaker than sequential consistency.

Globally:

6. **STORE**'s to the same location respect a **global total order**.
7. **LOCK** instructions respect a **global total order**.
8. Memory ordering preserves **transitive visibility** ("causality").

Impact of Reordering

Processor 0

1	mov 1, a	<i>;Store</i>
2	mov b, %ebx	<i>;Load</i>

Processor 1

3	mov 1, b	<i>;Store</i>
4	mov a, %eax	<i>;Load</i>

Impact of Reordering

Processor 0

1 mov 1, a ;Store
2 mov b, %ebx ;Load

2 mov b, %ebx ;Load
1 mov 1, a ;Store

Processor 1

3 mov 1, b ;Store
4 mov a, %eax ;Load

4 mov a, %eax ;Load
3 mov 1, b ;Store

The ordering $\langle 2, 4, 1, 3 \rangle$ produces $\text{\%eax} = \text{\%ebx} = 0$.

Instruction reordering violates sequential consistency!


Further Impact of Reordering

- The **LOAD**'s of **B_wants** and **A_wants** can be reordered before the **STORE**'s of **A_wants** and **B_wants**, respectively
- Both **Alice** and **Bob** might enter their critical sections simultaneously!

Peterson's algorithm revisited


Alice

```
A_wants = true;  
turn = B;  
while (B_wants && turn==B);  
frob(&x); //critical section  
A_wants = false;
```



Bob

```
B_wants = true;  
turn = A;  
while (A_wants && turn==A);  
borf(&x); //critical section  
B_wants = false;
```



Restoring Consistency

- Use `memory fence` to prevent hardware reordering
- Use `volatile` to prevent compiler reordering

Memory Fences

- A **memory fence** (or **memory barrier**) is a hardware action that enforces an **ordering** constraint between the instructions before and after the fence
- A memory fence can be done **explicitly** as an instruction (x86: **mfence**) or **implicitly** by locking, exchanging, and other synchronizing instructions
- The LLVM compiler implements a memory fence via the function **atomic_thread_fence()** defined in the C header file **stdatomic.h**.*
- The typical cost of a memory fence is comparable to that of an **L2-cache access**

*See <http://en.cppreference.com/w/c/atomic>.

Restoring Consistency

Alice

```
A_wants = true;
turn = B;
while (B_wants && turn==B);
frob(&x); //critical section
A_wants = false;
```

Bob

```
B_wants = true;
turn = A;
while (A_wants && turn==A);
borf(&x); //critical section
B_wants = false;
```

Memory fences can restore sequential consistency.

```
A_wants = true;
turn = B;
atomic_thread_fence();
while (B_wants && turn==B);
frob(&x); //critical section
A_wants = false;
```

```
B_wants = true;
turn = A;
atomic_thread_fence();
while (A_wants && turn==A);
borf(&x); //critical section
B_wants = false;
```

Well, sort of.

You also need to make sure that the **compiler** doesn't screw you over.

Restoring Consistency

- Memory fence to prevent hardware reordering
- Use `volatile` to prevent compiler reordering
 - ▣ It tells the compiler: “Do not optimize access to this variable. Always read/write it directly from memory.”

```
int data = 0;
bool ready = false;

void writer() {
    data = 42;      // Step 1
    ready = true;   // Step 2
}

void reader() {
    if (ready) {
        printf("Data: %d\n", data); // ⚠ Might print 0!
    }
}
```

```
int data = 0;
volatile bool ready = false; // 📌 Volatile flag

void writer() {
    data = 42;
    ready = true; // Write to volatile – prevents reordering
}

void reader() {
    if (ready) { // Read from volatile – re-fetch every time
        printf("Data: %d\n", data);
    }
}
```

So `reader()` might see `ready == true` and read `data == 0`, which is stale.

It **prevents reordering** of `data = 42` and `ready = true` at compile time.

Also ensures the reader thread **doesn't cache** `ready` in a register and miss updates.

Restoring Consistency with C11

```
widget x; //protected variable  
_Atomic bool A_wants = false;  
_Atomic bool B_wants = false;  
_Atomic enum {A, B} turn;
```

Alice

```
atomic_store(&A_wants, true);  
atomic_store(&turn, B);  
while (atomic_load(&B_wants) &&  
       atomic_load(&turn)==B);  
frob(&x); //critical section  
atomic_store(&A_wants, false);
```

Bob

```
atomic_store(&B_wants, true);  
atomic_store(&turn, A);  
while (atomic_load(&A_wants) &&  
       atomic_load(&turn)==A);  
borf(&x); //critical section  
atomic_store(&B_wants, false);
```

- The C11 language standard defines its own weak memory model, in which you can control hardware and compiler reordering of memory operations by
- declaring variables as `_Atomic`; and
- using the functions `atomic_load()`, `atomic_store()`, etc. as needed

See <http://en.cppreference.com/w/c/atomic>

Implementing General Mutexes

- **Theorem** [Burns-Lynch]. Any n -thread deadlock-free mutual-exclusion algorithm using only **LOAD** and **STORE** memory operations requires $\Omega(n)$ space
- **Theorem** [Attiya et al.]: Any n -thread deadlock-free mutual-exclusion algorithm on a modern machine must use an expensive operation such as a **memory fence** or an **atomic COMPARE-AND-SWAP** operation

Thus, hardware designers are justified when they implement special operations to support atomicity.

Outline

- Sequential Consistency (SC)
 - Mutual Exclusion w/o Locks under SC
- Relaxed Memory Consistency
 - Total Store Order (TSO)
- **Compare and Swap (CAS)**
- Lock-free Algorithms
 - Lock-free Stack
 - The ABA Problem

The Lock-Free Toolbox

Memory operations

- LOAD
- STORE
- CAS (COMPARE-AND-SWAP)



Compare-and-Swap

Specification

```
bool CAS(T *x, T old, T new) {  
    if (*x == old) {  
        *x = new;  
        return true;  
    }  
    return false;  
}
```

- Executes atomically.
- Implicit fence.

- COMPARE-AND-SWAP is provided by the `cmpxchg` instruction on x86-64
- The C header file `stdatomic.h` provides `CAS` via the built-in function `atomic_compare_exchange_strong()` which can operate on various integer types.*

* See http://en.cppreference.com/w/cpp/atomic/atomic_compare_exchange.

Mutex Using CAS

- **Theorem.** An n -thread deadlock-free mutual-exclusion algorithm using **CAS** can be implemented using $\Theta(1)$

Proof.

```
void lock(int *lock_var) {  
    while (!CAS(*lock_var, false, true));  
}
```

Is it free?

```
void unlock(int *lock_var) {  
    *lock_var = false;  
}
```

Yes?
Acquire it

Just the space for the mutex itself. ■

release it

Mutex Using CAS

- **Theorem.** An n -thread deadlock-free mutual-exclusion algorithm using **CAS** can be implemented using $\Theta(1)$ space.

Theorem [Burns-Lynch]. Any n -thread deadlock-free mutual-exclusion algorithm using only **LOAD** and **STORE** memory operations requires $\Omega(n)$ space

Summing Problem

```
int compute(const X& v);
int main() {
    const int n = 1000000;
    extern X myArray[n];
    // ...

    int result = 0;
    cilk_for (int i = 0; i < n; ++i) {
        result += compute(myArray[i]);
    }
    printf("The result is: %f\n", result );
    return 0;
}
```

Race!

Mutex Solution

```
int compute(const X& v);
int main() {
    const int n = 1000000;
    extern X myArray[n];
    mutex L;
    // ...

    int result = 0;
    cilk_for (int i = 0; i < n; ++i) {
        int temp = compute(myArray[i]);
        L.lock();
        result += temp;
        L.unlock();
    }
    printf( "The result is: %f\n", result );
    return 0;
}
```

Mutex Solution

Yet all we want is to atomically execute a **LOAD** of **x** followed by a store of **x**.

```
int compute(const X& v);
int main() {
    const int n = 1000000;
    extern X myArray[n];
    mutex L;
    // ...

    int result = 0;
    cilk_for (int i = 0; i < n; ++i) {
        int temp = compute(myArray[i]);
        L.lock();
        result += temp;
        L.unlock();
    }
    printf( "The result is: %f\n", result );
    return 0;
}
```

Q. What happens if the OS swaps out a loop iteration just after it acquires the mutex?

A. All other loop iterations must wait.

CAS Solution

```
int result = 0;
cilk_for (int i = 0; i < n; ++i) {
    int temp = compute(myArray[i]);
    int old, new;
    do {
        old = result;
        new = old + temp;
    } while (!CAS(&result, old, new));
}
```

Q. Now what happens if the OS swaps out a loop iteration?


A. No other loop iteration needs to wait. The algorithm is **nonblocking**

Outline

- Sequential Consistency (SC)
 - Mutual Exclusion w/o Locks under SC
- Relaxed Memory Consistency
 - Total Store Order (TSO)
- Compare and Swap (CAS)
- **Lock-free Algorithms**
 - Lock-free Stack
 - The ABA Problem


Lock-Free Stack

```
struct Node {  
    Node* next;  
    int data;  
};  
  
struct Stack {  
    Node* head;  
    :  
};
```




Lock-Free PUSH

```
void push(Node* node) {  
    do {  
        node->next = head;  
    } while (!CAS(&head, node->next, node));  
}
```



Lock-Free PUSH with Contention


```
void push(Node* node) {  
    do {  
        node->next = head;  
    } while (!CAS(&head, node->next, node));  
}
```



The COMPARE-AND-SWAP fails!

Lock-Free POP

```
Node* pop() {  
    Node* current = head;  
    while (current) {  
        if (CAS(&head, current, current->next)) break;  
        current = head;  
    }  
    return current;  
}
```



Performance Considerations

- COMPARE-AND-SWAP acquires a cache line in exclusive mode, invalidating the cache line in other caches.
- **Result:** High contention if all processors **CAS** to same cache line.

Better: check before CAS

- First, read the memory location to check whether the value changed before attempting a **CAS**.
- Only **CAS** if the value didn't change.

Similar to the **test and spin lock** that reads the lock status before attempting the **xchg** operation.

Lock-Free Push and Pop

```
void push(Node* node) {  
    do {  
        node->next = head;  
    } while (head != node->next ||  
             !CAS(&head, node->next, node));  
}
```

```
Node* pop() {  
    Node* current = head;  
    while (current) {  
        if (head == current &&  
            CAS(&head, current, current->next)) break;  
        current = head;  
    }  
    return current;  
}
```

Even better: add exponential backoff

Lock-Free Data Structures

- Efficient lock-free algorithms are known for a variety of data structures
 - e.g., linked lists, queues, skip lists, hash tables
- In theory, a thread might **starve**. In practice, starvation **rarely** happens
 - Because of contention, its operation might never complete
- **Transactional memory** offers another way to mutual exclusion w/o locks
 - TM allows a block of code to execute atomically


Practical issues with lock-free programming

- memory management
- contention
- the ABA problem

Outline


- Sequential Consistency (SC)
 - Mutual Exclusion w/o Locks under SC
- Relaxed Memory Consistency
 - Total Store Order (TSO)
- Compare and Swap (CAS)
- Lock-free Algorithms
 - Lock-free Stack
 - **The ABA Problem**

ABA Example




1. Strand 1 begins to pop the node containing **15**, but stalls after reading **current->next**.

ABA Example




1. Strand 1 begins to pop the node containing **15**, but stalls after reading **current->next**.
2. Strand 2 pops the node containing **15**.

ABA Example




1. Strand 1 begins to pop the node containing **15**, but stalls after reading **current->next**.
2. Strand 2 pops the node containing **15**.
3. Strand 2 pops the node containing **94**.

ABA Example




1. Strand 1 begins to pop the node containing **15**, but stalls after reading **current->next**.
2. Strand 2 pops the node containing **15**.
3. Strand 2 pops the node containing **94**.
4. Strand 2 pushes the value **7**, reusing the node that contained **15**.

ABA Example



1. Strand 1 begins to pop the node containing **15**, but stalls after reading **current->next**.
2. Strand 2 pops the node containing **15**.
3. Strand 2 pops the node containing **94**.
4. Strand 2 pushes the node **7**, reusing the node that contained **15**.
5. Strand 1 resumes, and its **CAS** succeeds, removing **7**, but putting **garbage** back on the list.

ABA Example



So we went from the head being

1. **A** = node with 15
2. To **B** = node with 94
3. Back to **A** = node that had 15 (but now has 7)

Looks the same but the world changed completely in the interim...

Notice that checking the value in the node will not help since in step 3 we could have inserted 15 back and gotten the same error...

Solutions to ABA

Versioning

- Pack a **version #** with each pointer in the same atomically updatable word
- Increment the version number every time the pointer is changed
- **CAS** both the pointer and the **version #** as a single atomic operation

Issue

- Version numbers may need to be very large.

Reclamation

- Prevent node reuse while pending requests exist.
- e.g. prevent node **15** from being reused as **7** while Strand 1 still executing

Summary

- Sequential Consistency
 - Peterson's algorithm for mutual exclusion
- Relaxed Memory Consistency
 - TSO: total store order (performance vs. productivity)
 - Use `memory fence` to prevent hardware reordering
 - Use `volatile` to prevent compiler reordering
- Compare-and-Swap (CAS)
- Use CAS to implement lock-free algorithms