



Software
Performance
Engineering

LECTURE 5

Vectorization

Xuhao Chen

September 29, 2026





Previously

- Performance is the currency of computing
- Bit hacks for **word** level parallelism
- Bentley rules to reduce work
- Understand your HW: Assembly, ISA, **instruction** level parallelism
- **Today:** **data** parallelism

Vectorization


- The First Supercomputers were developed in the 70's
- 200,000 gates, 80 MHz, \$10M



Seymour Cray

Vectorization

- The First Supercomputers were developed in the 70's
- Required long vectors for performance



Seymour Cray

Multimedia Instructions


- Multimedia ISA extensions came in the 90's
 - **SIMD: Single Instruction Multiple Data**
 - Integrated into the processor
 - Short and fast
 - No startup overhead

Vendor	ISA	Year	Bitwidth
Intel	MMX	1996	64 bits
AMD	3DNow!	1998	64 bits
Intel	SSE	1999	128 bits
Intel	SSE2	2001	128 bits
Intel	SSE3	2006	128 bits
Intel	AVX	2008	256 bits
ARM	Neon	2011	128 bits
Intel	AVX2	2013	256 bits
Intel	AVX512	2015	512 bits
Intel	AVX512-VNNI	2021	512 bits

Outline

- **What is SIMD**
- **X86 SIMD Instructions**
- **Auto-vectorization (by compiler)**
- **Hand Vectorization using Intrinsics**



Flynn Taxonomy (1966)



Flynn Taxonomy (1966)

Single-Instruction Single-Data



(e.g. Single-Core Processors)



Flynn Taxonomy (1966)

Single-Instruction Single-Data



(e.g. Single-Core Processors)



Flynn Taxonomy (1966)

Single-Instruction Single-Data

(e.g. Single-Core Processors)



LOAD




ADD 5



MUL 2



STORE



LOAD



ADD 5



MUL 2





STORE




Single-Instruction Multi-Data


(e.g. GPUs, Intel SIMD Extensions)



VLOAD




VADD 5




VMUL 2



VSTORE



VLOAD





VADD 5



VMUL 2




VSTORE



Flynn Taxonomy (1966)


Single-Instruction Single-Data

(e.g. Single-Core Processors)




Single-Instruction Multi-Data

(e.g. GPUs, Intel SIMD Extensions)




Multi-Instruction Single-Data

(e.g. Systolic Arrays,...)




Multi-Instruction Multi-Data

(e.g. Parallel Computers)




Scalar Hardware

Modern microprocessors spend a lot of real estate in **instruction processing**.
The data paths and ALUs are relatively small.




Vector Hardware

Modern microprocessors incorporate **vector units** to process data in a **single-instruction stream, multiple-data stream (SIMD)** fashion




A k-Wide Vector Unit



Each **vector register** holds k scalar integer or floating-point values

The unit contains k **vector lanes**, each containing scalar integer or floating-point hardware

A k-Wide Vector Unit




All vector lanes operate in **lock-step** and use the **same** instruction and control signals

Lock step!


SIMD Instructions

- SIMD instructions typically operate in an **elementwise** fashion
- The **i**th element of one vector register only takes part in operations with the **i**th element of other vector registers

(a) Scalar Operation




(b) SIMD Operation




SIMD Instructions

- SIMD instructions typically operate in an **elementwise** fashion
- All lanes perform **exactly the same operation** on each element



Example of SIMD
Processable Patterns





Example of SIMD
Unprocessable
Patterns

SIMD Instructions


- SIMD instructions typically operate in an **elementwise** fashion
- All lanes perform **exactly the same operation** on each element

Vector Addition




SIMD Instructions

- SIMD instructions typically operate in an **elementwise** fashion
- Vector memory operands might need to be **aligned**




4-byte memory access for aligned data



4-byte memory access for misaligned data


SIMD Instructions

- SIMD instructions typically operate in an **elementwise** fashion
- Some architectures support **cross-lane** operations
 - e.g. **inserting** or **extracting** subsets of vector elements
 - e.g. **permuting** (a.k.a., **shuffling**) the vector
 - e.g. **scatter**, or **gather**




Vector Register Aliasing

- Different names, but same registers



SIMD Data Types

- Each AVX register can have different data types
- Operation on 32 8-bit values in one instruction!



Outline


- What is SIMD
- **X86 SIMD Instructions**
- Auto-vectorization (by compiler)
- Hand Vectorization using Intrinsics

Instruction Naming Convention

- Format: `_mm<vec-width>_<op>_<scalar-ty>`
- `<vec-width>`: 64 / 128 / 256 / 512
- `<scalar-ty>`
 - **ss** – one single precision floating point (scalar)
 - **sd** – one double precision floating point (scalar)
 - **ps** - packed single precision floating point (vector)
 - **pd** - packed double precision floating point (vector)
 - **epi8/epi16/epi32/epi64** - packed signed integers (vector)


Instruction Naming Convention

- 256 bit 4-wide double precision float addition



Instruction Naming Convention

- 128-bit 4-wide load (128 is omitted) for single precision float



Instruction Naming Convention

- More Examples
 - 8-wide float addition: `_mm256_add_ps`
 - 16-wide absolute value of 16-bit ints: `_mm256_abs_epi16`
 - 4-wide float subtraction: `_mm_sub_ps` (128 is omitted)

Categories of Vector Instructions

1. Arithmetic and Logic
2. Compare
3. Load/Store
4. Shuffle (Swizzle)
5. Masks
6. Non-SIMD


1. Arithmetic and Logic

- **Types**

- add, sub, mul, div, and, or, not, xor

- **Example**

```
__m128i _mm_add_epi32 (__m128i a, __m128i b) {  
    for j := 0 to 3 {  
        i := j*32  
        dst[i+31:i] := a[i+31:i] + b[i+31:i]  
    }  
    return dst  
}
```




1. Arithmetic and Logic

- **Types**

- mul

- **Example**

```
__m256i _mm_mul_epu32 (__m256i a, __m256i b) {  
    for j := 0 to 3 {  
        i := j*64  
        dst[i+63:i] := a[i+31:i] * b[i+31:i]  
    }  
    dst[MAX:256] := 0  
    return dst  
}
```



2. Compare

- **Types**

- "Normal" comparisons: sets lanes to ones if true
- AVX-512 comparisons: sets mask register (%k0 - %k7)

- **Example**

```
__m128 _mm_cmp_ps (__m128 a, __m128 b, const int imm8) {  
    CASE (imm8[4:0]) {  
        0: OP := _CMP_EQ_OQ  
        1: OP := _CMP_LT_OS  
        ...  
        31: OP := _CMP_TRUE_US  
    }  
    for j := 0 to 3 {  
        i := j*32  
        dst[i+31:i] := ( a[i+31:i] OP b[i+31:i] ) ? 0xFFFFFFFF : 0  
    }  
    dst[MAX:128] := 0  
}
```

3. Load/Store

- **Types**

- Aligned load/store (e.g., vmovaps)
- Unaligned load/store (e.g., vmovups)
- Streaming load/store (e.g., vmovntps)
 - Streaming writes are 2x faster than normal writes


- **Example**

```
__m256 _mm256_loadu_ps (float const * mem_addr) {  
    dst[255:0] := MEM[mem_addr+255:mem_addr]  
    dst[MAX:256] := 0  
}
```

4. Shuffle: Blending


- Shuffle operation: `blend`
- **Example**

`__m128 _mm_blendv_ps(__m128 v1, __m128 v2, __m128 mask)`



4. Shuffle: all-to-all shuffle

- Shuffle operation: `shuffle`
- **Example**




each byte of `xmm2/m128`:
bit 7 == 0 specifies copying, bit 3:0 specifies 0 to 15
bit 7 == 1 specifies zero clearing

4. Shuffle: unpack

- Shuffle operation: unpackhi & unpacklo
- **Example**


`__m128 _mm_unpackhi_ps(__m128 a, __m128 b)`



5. Non-SIMD: dot product instructions

- Non-SIMD operation: madd
- **Example**

`__m128i _mm_maddubs_epi16(__m128i a, __m128i b)`



Calculate $A*B+C*D$ and set to each WORD of (3). Signed saturation: set -32768 / 32767 on overflow. Each BYTE of (1) is unsigned. Each BYTE of (2) is signed.


5. Non-SIMD: addsub

- Non-SIMD operation: addsub

$$(a+bi)(c+di) = (ac - bd) + (ad + bc)i$$

- Example**

```
__m256d _mm256_addsub_pd(__m256d a, __m256d b) {  
  for j := 0 to 3 {  
    i := j*64  
    if ((j & 1) == 0)  
      dst[i+63:i] := a[i+63:i] - b[i+63:i]  
    else  
      dst[i+63:i] := a[i+63:i] + b[i+63:i]  
  }  
}
```



And a lot more instructions..

Intel® Intrinsics Guide

Updated
08/10/2022

Version
3.6.3

Instruction Set

- ☐ MMX
- ☐ SSE
- ☐ SSE2
- ☐ SSE3
- ☐ SSE3
- ☐ SSE4.1
- ☐ SSE4.2
- ☐ AVX
- ☐ AVX2
- ☐ FMA
- ☐ AVX_VNNI
- ☐ AVX-512
- ☐ KNC
- ☐ AMX
- ☐ SVM
- ☐ Other

Categories

- ☐ Application-Targeted
- ☐ Arithmetic
- ☐ Bit Manipulation
- ☐ Cast
- ☐ Compare
- ☐ Convert
- ☐ Cryptography
- ☐ Elementary Math Functions
- ☐ General Support
- ☐ Load
- ☐ Logical
- ☐ Mask
- ☐ Miscellaneous
- ☐ Move
- ☐ OS-Targeted
- ☐ Probability/Statistics
- ☐ Random
- ☐ Set
- ☐ Shift

Search Intel Intrinsics

```
void __mm2intersect_epi32 (__m128i a, __m128i b, __mmask8* k1, __mmask8* k2) vp2intersectd
void __mm256_2intersect_epi32 (__m256i a, __m256i b, __mmask8* k1, __mmask8* k2) vp2intersectd
void __mm512_2intersect_epi32 (__m512i a, __m512i b, __mmask16* k1, __mmask16* k2) vp2intersectd
void __mm2intersect_epi64 (__m128i a, __m128i b, __mmask8* k1, __mmask8* k2) vp2intersectq
void __mm256_2intersect_epi64 (__m256i a, __m256i b, __mmask8* k1, __mmask8* k2) vp2intersectq
void __mm512_2intersect_epi64 (__m512i a, __m512i b, __mmask8* k1, __mmask8* k2) vp2intersectq
__m512i __mm512_4dpwssd_epi32 (__m512i src, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b) vp4dpwssd
__m512i __mm512_mask_4dpwssd_epi32 (__m512i src, __mmask16 k, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b) vp4dpwssd
__m512i __mm512_maskz_4dpwssd_epi32 (__mmask16 k, __m512i src, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b) vp4dpwssd
__m512i __mm512_4dpwssds_epi32 (__m512i src, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b) vp4dpwssds
__m512i __mm512_mask_4dpwssds_epi32 (__m512i src, __mmask16 k, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b) vp4dpwssds
__m512i __mm512_maskz_4dpwssds_epi32 (__mmask16 k, __m512i src, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b) vp4dpwssds
__m512 __mm512_4fmadd_ps (__m512 src, __m512 a0, __m512 a1, __m512 a2, __m512 a3, __m128 * b) v4fmaddps
__m512 __mm512_mask_4fmadd_ps (__m512 src, __mmask16 k, __m512 a0, __m512 a1, __m512 a2, __m512 a3, __m128 * b) v4fmaddps
__m512 __mm512_maskz_4fmadd_ps (__mmask16 k, __m512 src, __m512 a0, __m512 a1, __m512 a2, __m512 a3, __m128 * b) v4fmaddps
__m128 __mm_4fmadd_ss (__m128 src, __m128 a0, __m128 a1, __m128 a2, __m128 a3, __m128 * b) v4fmaddss
__m128 __mm_mask_4fmadd_ss (__m128 src, __mmask8 k, __m128 a0, __m128 a1, __m128 a2, __m128 a3, __m128 * b) v4fmaddss
```

Synopsis

```
__m128 __mm_mask_4fmadd_ss (__m128 src, __mmask8 k, __m128 a0, __m128 a1, __m128 a2, __m128 a3, __m128 * b)
#include <immintrin.h>
Instruction: v4fmaddss xmm (k), xmm, m128
CPUID Flags: AVX512_4FMAPS
```

Description

Multiply the lower single-precision (32-bit) floating-point elements specified in 4 consecutive operands `a0` through `a3` by corresponding element in `b`, accumulate with the lower element in `a`, and store the result in the lower element of `dst` using writemask `k` (the element is copied from `a` when mask bit 0 is not set).

Operation

```
dst[127:0] := src[127:0]
IF k[0]
    FOR m := 0 to 3
        addr := b + m * 32
        dst.fp32[0] := dst.fp32[0] + a[m].fp32[0] * Cast_FP32(MEM[addr+31:addr])
    ENDFOR
FI
dst[MAX:128] := 0
```

```
__m128 __mm_maskz_4fmadd_ss (__mmask8 k, __m128 src, __m128 a0, __m128 a1, __m128 a2, __m128 a3, __m128 * b) v4fmaddss
__m512 __mm512_4fmadd_ps / __m512 src, __m512 a0, __m512 a1, __m512 a2, __m512 a3, __m128 * b) v4fmaddps
```

> 7000 Intrinsics!

Give Feedback

Outline

- What is SIMD
- X86 SIMD Instructions
- **Auto-vectorization (by compiler)**
- Hand Vectorization using Intrinsics

Vectorizing Compiler

- Two types of Vectorizers
 - Loop Vectorizer
 - Superword Level Parallelism (SLP) Vectorizer

Loop Vectorization

```
for (int i = 0; i < n; i++) {  
    tb  = b[i];  
    tc  = c[i];  
    t   = tb + tc;  
    a[i] = t;  
}
```

Loop Vectorization

```
for (int i = 0; i < n; i++) {  
    tb  = b[i];  
    tc  = c[i];  
    t   = tb + tc;  
    a[i] = t;  
}
```

```
for (int i = 0; i < n; i += 2) {  
    tb  = b[i:i+2];  
    tc  = c[i:i+2];  
    t   = tb + tc;  
    a[i:i+2] = t;  
}
```

The entire loop body is vectorized.

In order to do this the loop has to be ‘parallelizable’

Non-trivial analysis by the compiler

SLP Vectorization

```
for (int i = 0; i < n; i++) {  
    tb  = b[i];  
    tc  = c[i];  
    t   = tb + tc;  
    a[i] = t;  
}
```

```
for (int i = 0; i < n; i += 2) {  
    tb  = b[i];  
    tc  = c[i];  
    t   = tb + tc;  
    a[i] = t;  
    tb2 = b[i+1];  
    tc2 = c[i+1];  
    t2  = tb2 + tc2;  
    a[i] = t2;  
}
```

First Unroll the loop

SLP Vectorization

```
for (int i = 0; i < n; i++) {  
    tb  = b[i];  
    tc  = c[i];  
    t   = tb + tc;  
    a[i] = t;  
}
```

```
for (int i = 0; i < n; i += 2) {  
    tb  = b[i];  
    tc  = c[i];  
    t   = tb + tc;  
    a[i] = t;  
    tb2 = b[i+1];  
    tc2 = c[i+1];  
    t2  = tb2 + tc2;  
    a[i] = t2;  
}
```

Next, group isomorphic instructions

SLP Vectorization

```
for (int i = 0; i < n; i++) {  
    tb  = b[i];  
    tc  = c[i];  
    t   = tb + tc;  
    a[i] = t;  
}
```

```
for (int i = 0; i < n; i += 2) {  
    tb  = b[i];  
    tb2 = b[i+1];  
    tc  = c[i];  
    tc2 = c[i+1];  
    t   = tb + tc;  
    t2  = tb2 + tc2;  
    a[i] = t;  
    a[i] = t2;  
}
```

Next, group isomorphic instructions

SLP Vectorization

```
for (int i = 0; i < n; i++) {  
    tb  = b[i];  
    tc  = c[i];  
    t   = tb + tc;  
    a[i] = t;  
}
```

```
for (int i = 0; i < n; i += 2) {  
    tb  = b[i];  
    tb2 = b[i+1];  
    tc  = c[i];  
    tc2 = c[i+1];  
    t   = tb + tc;  
    t2  = tb2 + tc2;  
    a[i] = t;  
    a[i+1] = t2;  
}
```

```
tb      = b[i:i+1];  
tc      = c[i:i+1];  
t       = tb + tc;  
a[i:i+1] = t;
```

Finally, pack into vector instructions

SLP Vectorization


```
for (int i = 0; i < n; i++) {  
    tb  = b[i];  
    tc  = c[i];  
    t   = tb + tc;  
    a[i] = t;  
}
```

```
tb      = b[i:i+1];  
tc      = c[i:i+1];  
t       = tb + tc;  
a[i:i+1] = t;
```

The loop body can be partially vectorized.
No loop analysis is needed, only the loop body.
Need to unroll to be effective

How good are the Vectorizers?

- Vectorizers are good, but, not ubiquitous
 - Some complex code get vectorized
 - But... some very simple code does not!
 - Need to pay attention



S. Maleki, Y. Gao, M. J. Garzarn, T. Wong and D. A. Padua, "An Evaluation of Vectorizing Compilers," 2011 International Conference on Parallel Architectures and Compilation Techniques, 2011

How good are the Vectorizers: A Case Study

- A matrix-vector multiplication kernel from TVM

```
void
dot_16x16x16_uint8_int8_int32 (
  uint8_t data[restrict 4],
  int8_t kernel[restrict 16][4],
  int32_t output[restrict 16]) {
  for (int i = 0; i < 16; i++)
    for (int k = 0; k < 4; k++)
      output[i] +=
        data[k] * kernel[i][k];
}
```

ICC (1x)

```
movzx r11d, [rdi]
movsx eax, [rsi]
imul r11d, eax
movsx ecx, [1+rsi]
movzx r8d, [1+rdi]
imul r8d, ecx
movzx r10d, [2+rdi]
movsx r9d, [2+rsi]
imul r10d, r9d
movsx eax, [3+rsi]
movzx ecx, [3+rdi]
imul ecx, eax
add r11d, [rdx]
add r11d, r8d
add r11d, r10d
add r11d, ecx
mov [rdx], r11d
movzx r9d, [rdi]
movsx r11d, [4+rsi]
imul r9d, r11d
movsx eax, [5+rsi]
movzx r8d, [1+rdi]
imul r8d, eax
movzx eax, [2+rdi]
movsx ecx, [6+rsi]
imul ecx, eax
movsx r10d, [7+rsi]
movzx r11d, [3+rdi]
imul r11d, r10d
add r9d, [4+rdx]
add r9d, r8d
add r9d, eax
add r9d, r11d
mov [4+rdx], r9d
movzx r8d, [rdi]
movsx r9d, [8+rsi]
imul r8d, r9d
movsx eax, [9+rsi]
movzx ecx, [1+rdi]
imul ecx, eax
movzx r10d, [2+rdi]
movsx r9d, [10+rsi]
imul r10d, r9d
movsx r11d, [11+rsi]
movzx eax, [3+rdi]
imul eax, r11d
add r8d, [8+rdx]
add r8d, ecx
add r8d, r10d
add r8d, eax
mov [8+rdx], r8d
movzx ecx, [rdi]
movsx r8d, [12+rsi]
```

GCC (1.5x)

```
vmovdqa xmm0, .LC0[rip]
vmovdqu xmm1, [rsi]
vmovdqu xmm3, [rsi+16]
vmovdqu xmm2, [rsi+32]
vmovdqu xmm6, [rsi+48]
vpand xmm10, xmm0, xmm1
vpsrlw xmm1, xmm1, 8
movzx r9d, [rdi]
vpand xmm4, xmm0, xmm3
vpsrlw xmm3, xmm3, 8
movzx ecx, [rdi+2]
movzx eax, [rdi+3]
vpacsubb xmm4, xmm10, xmm4
vpacsubb xmm1, xmm1, xmm3
vpand xmm10, xmm0, xmm2
movzx r8d, [rdi+1]
vpand xmm3, xmm0, xmm6
vpsrlw xmm2, xmm2, 8
vmovd xmm8, r9d
vpacsubb xmm3, xmm10, xmm3
vpsrlw xmm6, xmm6, 8
vpand xmm10, xmm0, xmm4
vpacsubb xmm2, xmm2, xmm6
vpsrlw xmm4, xmm4, 8
vpand xmm6, xmm0, xmm3
vpsrlw xmm3, xmm3, 8
vpacsubb xmm10, xmm10, xmm6
vpacsubb xmm8, xmm8, xmm8
vpacsubb xmm6, xmm4, xmm3
vpand xmm3, xmm0, xmm1
vpand xmm0, xmm0, xmm2
vpacsubb xmm3, xmm3, xmm0
vpsrldq xmm0, xmm10, 8
vpmovsxbw xmm4, xmm10
vmovd xmm7, r8d
vpmullw xmm4, xmm4, xmm8
vpmullw xmm0, xmm0, xmm8
vpsrlw xmm2, xmm2, 8
vpbroadcastw xmm7, xmm7
vpmovsxbw xmm8, xmm3
vpsrldq xmm3, xmm3, 8
vmovd xmm9, ecx
vpmullw xmm8, xmm8, xmm7
vpsrlw xmm1, xmm1, 8
vpbroadcastw xmm9, xmm9
vpmovsxbw xmm3, xmm3
vpacsubb xmm1, xmm1, xmm2
vmovd xmm5, eax
vpmullw xmm3, xmm3, xmm7
vpmovsxbw xmm7, xmm6
vpbroadcastw xmm5, xmm5
vpmullw xmm7, xmm7, xmm9
vpsrldq xmm2, xmm6, 8
vpacsubb xmm6, xmm6, xmm1
```

Clang/LLVM (2.2x)

```
movzx eax, [rdi + 3]
vpbroadcastd zmm8, eax
movzx eax, [rdi + 2]
vpbroadcastd zmm9, eax
movzx eax, [rdi + 1]
vpbroadcastd zmm10, eax
movzx eax, [rdi]
vpbroadcastd zmm11, eax
vmovdqu xmm0, [rsi]
vmovdqu xmm1, [rsi + 16]
vmovdqu xmm6, [rsi + 32]
vmovdqu xmm7, [rsi + 48]
vmovdqa xmm2, [rip + .LCPI0_0]
vpshufb xmm3, xmm7, xmm2
vpshufb xmm2, xmm6, xmm2
vpunpckldq xmm2, xmm2, xmm3
vmovdqa xmm3, [rip + .LCPI0_1]
vpshufb xmm4, xmm1, xmm3
vpshufb xmm3, xmm0, xmm3
vpunpckldq xmm3, xmm3, xmm4
vpblendd xmm12, xmm3, xmm2, 12
vmovdqa xmm3, [rip + .LCPI0_2]
vpshufb xmm4, xmm7, xmm3
vpshufb xmm3, xmm6, xmm3
vpunpckldq xmm3, xmm3, xmm4
vmovdqa xmm4, [rip + .LCPI0_3]
vpshufb xmm5, xmm1, xmm4
vpshufb xmm4, xmm0, xmm4
vpunpckldq xmm4, xmm4, xmm5
vpblendd xmm3, xmm4, xmm3, 12
vmovdqa xmm4, [rip + .LCPI0_4]
vpshufb xmm5, xmm7, xmm4
vpshufb xmm4, xmm6, xmm4
vpunpckldq xmm4, xmm4, xmm5
vmovdqa xmm5, [rip + .LCPI0_5]
vpshufb xmm2, xmm1, xmm5
vpshufb xmm5, xmm0, xmm5
vpunpckldq xmm2, xmm5, xmm2
vpblendd xmm2, xmm2, xmm4, 12
vmovdqa xmm4, [rip + .LCPI0_6]
vpshufb xmm5, xmm7, xmm4
vpshufb xmm4, xmm6, xmm4
vpunpckldq xmm4, xmm4, xmm5
vmovdqa xmm5, [rip + .LCPI0_7]
vpshufb xmm1, xmm1, xmm5
vpshufb xmm0, xmm0, xmm5
vpunpckldq xmm0, xmm0, xmm1
vpblendd xmm0, xmm0, xmm4, 12
vpmovsxbd zmm1, xmm12
vpmulld zmm1, zmm11, zmm1
vpadd zmm1, zmm1, [rdx]
vpmovsxbd zmm3, xmm3
vpmulld zmm3, zmm10, zmm3
vpmovsxbd zmm2, xmm2
```

VeGen (11x)

```
vmovdqu64 zmm0, [rdx]
vpbroadcastd zmm1, [rdi]
vpdpbusd zmm0, zmm0, [rsi]
vmovdqu64 [rdx], zmm0
```

What makes it hard for the compiler to vectorize

1	Loop with unknown trip count
2	Cannot prove independence
3	Loop carried dependences
4	Reductions
5	Control-flow
6	Non-contiguous data
7	Cost model failures
8	Vectorization of function calls
9	Different hardware versions

See <https://llvm.org/docs/Vectorizers.html> for more info.

What makes it hard for the compiler to vectorize

Go to godbolt.org and try

`-O3 -mavx512vnni -ffast-math -fno-unroll-loops`

1	Loop with unknown trip count	https://www.godbolt.org/z/MacxsGWsv
2	Cannot prove independence	https://www.godbolt.org/z/8fYcPz3Wv
3	Loop carried dependences	https://www.godbolt.org/z/z146WG3de
4	Reductions	https://www.godbolt.org/z/4n8zEbbsT
5	Control-flow	https://www.godbolt.org/z/jojKnEnv4
6	Non-contiguous data	https://www.godbolt.org/z/PzYc9jqaf https://www.godbolt.org/z/ME15GTesM
7	Cost model failures	https://godbolt.org/z/r5szGsaK4
8	Vectorization of function calls	https://www.godbolt.org/z/5Mcq9Pors
9	Different hardware versions	https://www.godbolt.org/z/Pvn1KM3rx

See <https://llvm.org/docs/Vectorizers.html> for more info.

1. Loop with unknown trip count

```
void bar(float * A, float K, int start, int end) {  
    for (int i = start; i < end; ++i)  
        A[i] = K;  
}
```

```
void bar(float * A, float K, int n) {  
    for (int i = 0; i < n; ++i)  
        A[i] = K;  
}
```

```
void bar(float * A, float K) {  
    for (int i = 0; i < 1024; ++i)  
        A[i] = K;  
}
```

1. Loop with unknown trip count

<pre> 1 bar: # @bar 2 cmpl %edx, %esi 3 jge .LBB0_8 4 movslq %esi, %rsi 5 movslq %edx, %rax 6 movq %rax, %rcx 7 subq %rsi, %rcx 8 cmpq \$8, %rcx 9 jae .LBB0_3 10 movq %rsi, %r8 11 jmp .LBB0_14 12 .LBB0_3: 13 cmpq \$16, %rcx 14 jae .LBB0_9 15 xorl %edx, %edx 16 jmp .LBB0_5 17 .LBB0_9: 18 movq %rcx, %rdx 19 andq \$-16, %rdx 20 vbroadcastss %xmm0, %zmm1 21 leaq (%rdi,%rsi,4), %r8 22 xorl %r9d, %r9d 23 .LBB0_10: # =>This Inner Loop Header: 24 vmovups %zmm1, (%r8,%r9,4) 25 addq \$16, %r9 26 cmpq %r8, %rdx 27 jne .LBB0_10 28 cmpq %rdx, %rcx 29 jle .LBB0_8 30 testb \$8, %cl 31 jne .LBB0_5 32 addq %rsi, %rdx 33 movq %rdx, %r8 34 jmp .LBB0_14 35 .LBB0_5: 36 movq %rcx, %r9 37 andq \$-8, %r9 38 leaq (%r9,%rsi), %r8 39 vbroadcastss %xmm0, %ymm1 40 leaq (%rdi,%rsi,4), %rsi 41 .LBB0_6: # =>This Inner Loop Header: 42 vmovups %ymm1, (%rsi,%rdx,4) 43 addq \$8, %rdx 44 cmpq %rdx, %r8 45 jne .LBB0_6 46 cmpq %r9, %rcx 47 je .LBB0_8 48 .LBB0_14: # =>This Inner Loop Header: 49 vmovss %xmm0, (%rdi,%r8,4) 50 incq %r8 51 cmpq %r8, %rax 52 jne .LBB0_14 53 .LBB0_8: 54 vzeroupper 55 retq </pre>	<pre> 1 bar: # @bar 2 testl %esi, %esi 3 jle .LBB0_13 4 movl %esi, %eax 5 cmpl \$8, %esi 6 jae .LBB0_3 7 xorl %ecx, %ecx 8 jmp .LBB0_12 9 .LBB0_3: 10 cmpl \$16, %esi 11 jae .LBB0_5 12 xorl %ecx, %ecx 13 jmp .LBB0_9 14 .LBB0_5: 15 movl %eax, %ecx 16 andl \$-16, %ecx 17 vbroadcastss %xmm0, %zmm1 18 leaq (,%rax,4), %rdx 19 andq \$-64, %rdx 20 xorl %esi, %esi 21 .LBB0_6: # =>This Inner Loop Header: Depth=1 22 vmovups %zmm1, (%rdi,%rsi) 23 addq \$64, %rsi 24 cmpq %rsi, %rdx 25 jne .LBB0_6 26 cmpq %rax, %rcx 27 je .LBB0_13 28 testb \$8, %al 29 je .LBB0_12 30 .LBB0_9: 31 movq %rcx, %rdx 32 movl %eax, %ecx 33 andl \$-8, %ecx 34 vbroadcastss %xmm0, %ymm1 35 .LBB0_10: # =>This Inner Loop Header: Depth=1 36 vmovups %ymm1, (%rdi,%rdx,4) 37 addq \$8, %rdx 38 cmpq %rdx, %rcx 39 jne .LBB0_10 40 cmpq %rax, %rcx 41 je .LBB0_13 42 .LBB0_12: # =>This Inner Loop Header: Depth=1 43 vmovss %xmm0, (%rdi,%rcx,4) 44 incq %rcx 45 cmpq %rcx, %rax 46 jne .LBB0_12 47 .LBB0_13: 48 vzeroupper 49 retq </pre>	<pre> 1 bar: # @bar 2 vbroadcastss %xmm0, %zmm0 3 xorl %eax, %eax 4 .LBB0_1: # =>This Inner Loop Header: Depth=1 5 vmovups %zmm0, (%rdi,%rax,4) 6 addq \$16, %rax 7 cmpq \$1024, %rax # imm = 0x400 8 jne .LBB0_1 9 vzeroupper 10 retq </pre>
---	--	---


2. Cannot Prove Independence

```
void bar(float * A, float * B) {  
    for (int i = 0; i < 4; ++i)  
        A[i] += B[i];  
}
```

1	2	3	4	5	6	7	8	10	20	30	40	50	60	70	80
---	---	---	---	---	---	---	---	----	----	----	----	----	----	----	----

2. Cannot Prove Independence

```
void bar(float * A, float * B) {  
    for (int i = 0; i < 4; ++i)  
        A[i] += B[i];  
}
```




1	2	3	4	5	6	7	8	10	20	30	40	50	60	70	80
---	---	---	---	---	---	---	---	----	----	----	----	----	----	----	----

Serial execution: in scalar

2. Cannot Prove Independence

```
void bar(float * A, float * B) {  
    for (int i = 0; i < 4; ++i)  
        A[i] += B[i];  
}
```




1	2	3	4	5	6	7	8	10	20	30	40	50	60	70	80
---	---	---	---	---	---	---	---	----	----	----	----	----	----	----	----

Parallel execution: in vector of width 2

2. Cannot Prove Independence

```
void bar(float * A, float * B) {  
    for (int i = 0; i < 4; ++i)  
        A[i] += B[i];  
}
```




1	2	3	4	5	6	7	8	10	20	30	40	50	60	70	80
---	---	---	---	---	---	---	---	----	----	----	----	----	----	----	----

Serial execution: in scalar

2. Cannot Prove Independence

- Vec
- Mer

```
void bar(float * A, float * B) {  
    for (int i = 0; i < 4; ++i)  
        A[i] += B[i];  
}
```



Parallel execution: in vector of width 2

2. Cannot Prove Independence

```
void bar(float * A, float * B) {  
    for (int i = 0; i < 1024; ++i)  
        A[i] += B[i];  
}
```

- **LLVM can still vectorize!**

2. Cannot Prove Independence

```
1  bar:                                     # @bar
2      leaq    4096(%rsi), %rax
3      cmpq    %rdi, %rax
4      jbe     .LBB0_2
5      leaq    4096(%rdi), %rax
6      cmpq    %rsi, %rax
7      jbe     .LBB0_2
8      xorl    %eax, %eax
9  .LBB0_6:                                # =>This Inner Loop Header: Depth=1
10     vmovss   (%rdi,%rax,4), %xmm0        # xmm0 = mem[0],zero,zero,zero
11     vaddss   (%rsi,%rax,4), %xmm0, %xmm0
12     vmovss   %xmm0, (%rdi,%rax,4)
13     incq     %rax
14     cmpq     $1024, %rax                 # imm = 0x400
15     jne     .LBB0_6
16     jmp     .LBB0_4
17 .LBB0_2:
18     xorl    %eax, %eax
19 .LBB0_3:                                # =>This Inner Loop Header: Depth=1
20     vmovups   (%rdi,%rax,4), %zmm0
21     vaddps   (%rsi,%rax,4), %zmm0, %zmm0
22     vmovups   %zmm0, (%rdi,%rax,4)
23     addq     $16, %rax
24     cmpq     $1024, %rax                 # imm = 0x400
25     jne     .LBB0_3
26 .LBB0_4:
27     vzeroupper
28     retq
```

sequential

parallel

2. Cannot Prove Independence

```
void bar(float * A, float * B) {  
    for (int i = 0; i < 1024; ++i)  
        A[i] += B[i];  
}
```

- LLVM can still vectorize!
- By checking if the address ranges overlap
 - Check the ranges first
 - If overlapping, use scalar implementation
 - If not overlapping, use the vectorized implementation
- However, Can save all these hassles by...

```
void bar(float * restrict A, float * restrict B) {  
    for (int i = 0; i < 1024; ++i)  
        A[i] += B[i];  
}
```

2. Cannot Prove Independence

```
1  bar:                                     # @bar
2      leaq    4096(%rsi), %rax
3      cmpq    %rdi, %rax
4      jbe     .LBB0_2
5      leaq    4096(%rdi), %rax
6      cmpq    %rsi, %rax
7      jbe     .LBB0_2
8      xorl    %eax, %eax
9  .LBB0_6:                                # =>This Inner Loop Header: Depth=1
10     vmovss   (%rdi,%rax,4), %xmm0        # xmm0 = mem[0],zero,zero,zero
11     vaddss   (%rsi,%rax,4), %xmm0, %xmm0
12     vmovss   %xmm0, (%rdi,%rax,4)
13     incq     %rax
14     cmpq     $1024, %rax                 # imm = 0x400
15     jne     .LBB0_6
16     jmp      .LBB0_4
17 .LBB0_2:
18     xorl     %eax, %eax
19 .LBB0_3:                                # =>This Inner Loop Header: Depth=1
20     vmovups   (%rdi,%rax,4), %zmm0
21     vaddps    (%rsi,%rax,4), %zmm0, %zmm0
22     vmovups   %zmm0, (%rdi,%rax,4)
23     addq      $16, %rax
24     cmpq      $1024, %rax                # imm = 0x400
25     jne      .LBB0_3
26 .LBB0_4:
27     vzeroupper
28     retq
```

Using restrict



```
bar:                                     # @bar
      xorl     %eax, %eax
.LBB0_1:                                # =>This Inner Loop Header
      vmovups   (%rdi,%rax,4), %zmm0
      vaddps    (%rsi,%rax,4), %zmm0, %zmm0
      vmovups   %zmm0, (%rdi,%rax,4)
      addq      $16, %rax
      cmpq      $1024, %rax                # imm = 0x400
      jne      .LBB0_1
      vzeroupper
      retq
```

3. Loop Carried Dependence

```
void bar(float * restrict A, int n) {  
    for (int i = 1; i < 1024; ++i)  
        A[i] += A[i-1];  
                A[i-2];  
                A[i-100];  
                A[i-128];  
}
```

```
void bar(float * restrict A, int n) {  
    for (int i = 0; i < n-1; ++i)  
        A[i] += A[i+1];  
}
```

3. Loop Carried Dependence

```
bar:                                     # @bar
    vmovss -4(%rdi), %xmm0              # xmm0 = mem[0],zero,zero,zero
    xorl   %eax, %eax
.LBB0_1:                                # =>This Inner Loop Header: Depth=1
    vaddss (%rdi,%rax,4), %xmm0, %xmm0
    vmovss %xmm0, (%rdi,%rax,4)
    incq   %rax
    cmpq   $1024, %rax                  # imm = 0x400
    jne    .LBB0_1
    retq
```

$A[i] += A[i-1];$

```
bar:                                     # @bar
    vmovsd -8(%rdi), %xmm0              # xmm0 = mem[0],zero
    xorl   %eax, %eax
.LBB0_1:                                # =>This Inner Loop Header: Depth=1
    vmovsd (%rdi,%rax,4), %xmm1         # xmm1 = mem[0],zero
    vaddps %xmm0, %xmm1, %xmm0
    vmovlps %xmm0, (%rdi,%rax,4)
    addq   $2, %rax
    cmpq   $1024, %rax                  # imm = 0x400
    jne    .LBB0_1
    retq
```

$A[i] += A[i-2];$

```
bar:                                     # @bar
    xorl   %eax, %eax
.LBB0_1:                                # =>This Inner Loop Header: Depth=1
    vmovups (%rdi,%rax,4), %xmm0
    vaddps -400(%rdi,%rax,4), %xmm0, %xmm0
    vmovups %xmm0, (%rdi,%rax,4)
    addq   $4, %rax
    cmpq   $1024, %rax                  # imm = 0x400
    jne    .LBB0_1
    retq
```


$A[i] += A[i-100];$

```
bar:                                     # @bar
    xorl   %eax, %eax
.LBB0_1:                                # =>This Inner Loop Header: Depth=1
    vmovups (%rdi,%rax,4), %xmm0
    vaddps -512(%rdi,%rax,4), %xmm0, %xmm0
    vmovups %xmm0, (%rdi,%rax,4)
    addq   $16, %rax
    cmpq   $1024, %rax                  # imm = 0x400
    jne    .LBB0_1
    vzeroupper
    retq
```

$A[i] += A[i-128];$

4. Reductions

```
int foo(int *A) {  
    unsigned sum = 0;  
    for (int i = 0; i < 1024; ++i)  
        sum += A[i] + 5;  
    return sum;  
}
```



4. Reductions

```
.LCPI0_0:
    .long    5                                # 0x5
foo:                                # @foo
    vpxor    %xmm0, %xmm0, %xmm0
    xorl     %eax, %eax
    vpbroadcastd .LCPI0_0(%rip), %zmm1    # zmm1 = [5,5,5,5,5,5,5,5,5,5,5,5,5,5,5,5]
.LBB0_1:                                # =>This Inner Loop Header: Depth=1
    vpaddd   (%rdi,%rax,4), %zmm0, %zmm0
    vpaddd   %zmm1, %zmm0, %zmm0
    addq     $16, %rax
    cmpq     $1024, %rax                    # imm = 0x400
    jne      .LBB0_1
    vextracti64x4 $1, %zmm0, %ymm1
    vpaddd   %zmm1, %zmm0, %zmm0
    vextracti128 $1, %ymm0, %xmm1
    vpaddd   %xmm1, %xmm0, %xmm0
    vpshufd  $238, %xmm0, %xmm1             # xmm1 = xmm0[2,3,2,3]
    vpaddd   %xmm1, %xmm0, %xmm0
    vpshufd  $85, %xmm0, %xmm1              # xmm1 = xmm0[1,1,1,1]
    vpaddd   %xmm1, %xmm0, %xmm0
    vmovd    %xmm0, %eax
    vzeroupper
    retq
```

broadcast

**16 elements at
a time**

**Reduction tree
on 16 elements**

5. Control Flow

```
int foo(int *A) {  
    unsigned sum = 0;  
    for (int i = 0; i < 1024; ++i)  
        if (A[i] > 0)  
            sum += A[i] + 5;  
    return sum;  
}
```

5. Control Flow

```
.LCPI0_0:
    .long    5                # 0x5
foo:
    # @foo
    vpxor    %xmm0, %xmm0, %xmm0
    xorl     %eax, %eax
    vpxor    %xmm1, %xmm1, %xmm1

.LBB0_1:
    # =>This Inner Loop Header: Depth=1
    vmovdqu64    (%rdi,%rax,4), %zmm2
    vpcmpgtd     %zmm0, %zmm2, %k1
    vpaddd      %zmm1, %zmm2, %zmm2
    vpaddd      .LCPI0_0(%rip){1to16}, %zmm2, %zmm1 {%k1}
    addq        $16, %rax
    cmpq        $1024, %rax          # imm = 0x400
    jne         .LBB0_1
    vextracti64x4    $1, %zmm1, %ymm0
    vpaddd      %zmm0, %zmm1, %zmm0
    vextracti128     $1, %ymm0, %xmm1
    vpaddd      %xmm1, %xmm0, %xmm0
    vpshufd     $238, %xmm0, %xmm1    # xmm1 = xmm0[2,3,2,3]
    vpaddd      %xmm1, %xmm0, %xmm0
    vpshufd     $85, %xmm0, %xmm1     # xmm1 = xmm0[1,1,1,1]
    vpaddd      %xmm1, %xmm0, %xmm0
    vmovd      %xmm0, %eax
    vzeroupper
    retq
```

add w/o
condition

Compute mask
 $A[i] > 0$?

Update value
only if mask set

6. Non-contiguous data

```
void bar(float * restrict A, float * restrict B, float K) {  
    for (int i = 0; i < 1024; ++i) {  
        B[2*i]    = A[i] + K;  
        B[2*i+1] = A[i] * K;  
    }  
}
```


6. Non-contiguous data

```
.LCPI0_0:
.long 0          # 0x0
.long 16         # 0x10
.long 1          # 0x1
.long 17         # 0x11
.long 2          # 0x2
.long 18         # 0x12
.long 3          # 0x3
.long 19         # 0x13
.long 4          # 0x4
.long 20         # 0x14
.long 5          # 0x5
.long 21         # 0x15
.long 6          # 0x6
.long 22         # 0x16
.long 7          # 0x7
.long 23         # 0x17

.LCPI0_1:
.long 8          # 0x8
.long 24         # 0x18
.long 9          # 0x9
.long 25         # 0x19
.long 10         # 0xa
.long 26         # 0x1a
.long 11         # 0xb
.long 27         # 0x1b
.long 12         # 0xc
.long 28         # 0x1c
.long 13         # 0xd
.long 29         # 0x1d
.long 14         # 0xe
.long 30         # 0x1e
.long 15         # 0xf
.long 31         # 0x1f

bar:
    vbroadcastss    %xmm0, %xmm0
    xorl    %eax, %eax
    vmovaps .LCPI0_0(%rip), %zmm1    # zmm1 = [0,16,1,17,2,18,3,19,4,20,5,21,6,22,7,23]
    vmovaps .LCPI0_1(%rip), %zmm2    # zmm2 = [8,24,9,25,10,26,11,27,12,28,13,29,14,30,15,31]

.LBB0_1:
    vmovups (%rdi,%rax,4), %zmm3
    vaddps %zmm0, %zmm3, %zmm4
    vmulps %zmm0, %zmm3, %zmm3
    vmovaps %zmm4, %zmm5
    vpermt2ps %zmm3, %zmm1, %zmm5
    vpermt2ps %zmm3, %zmm2, %zmm4
    vmovups %zmm4, 64(%rsi,%rax,8)
    vmovups %zmm5, (%rsi,%rax,8)
    addq $16, %rax
    cmpq $1024, %rax    # imm = 0x400
    jne .LBB0_1
    vzeroupper
    retq
```




6. Non-contiguous data

```
void bar(float * restrict A, float * restrict B, float K) {  
    for (int i = 0; i < 1024; ++i) {  
        B[2*i]    = A[i] + K;  
        B[2*i+1] = A[i] * K;  
    }  
}
```

```
void bar(float * restrict A, float * restrict B, int * restrict X) {  
    for (int i = 0; i < 1024; ++i)  
        B[i] *= A[X[i]];  
}
```

Indirect Access: Use gather instruction

6. Non-contiguous data



```
bar:                                     # @bar
    xorl    %eax, %eax
.LBB0_1:                                # =>This Inner Loop Header: Depth=1
    vmovups (%rdx,%rax,4), %zmm0
    kxnorw  %k0, %k0, %k1
    vxorps  %xmm1, %xmm1, %xmm1
    vgatherdps (%rdi,%zmm0,4), %zmm1 {%k1}
    vmovups %zmm1, (%rsi,%rax,4)
    addq    $16, %rax
    cmpq    $1024, %rax                  # imm = 0x400
    jne     .LBB0_1
    vzeroupper
    retq
```

gather from
memory

7. Cost Model Failures

- Clang v13 vs Clang v14
 - v13 is capable of vectorizing but thinks **not profitable**.
 - v14's updated cost model now thinks it's **profitable** to vectorize.
 - You can get about 30% speedup

```
void matmul(float a[restrict][N], float b[restrict][N],  
            float c[restrict][N], int N) {  
    for (int i = 0; i < N; i++)  
        for (int j = 0; j < N; j++)  
            for (int k = 0; k < N; k++)  
                c[i][j] += a[i][k] * b[k][j];  
}
```

7. Cost Model Failures

```

1  matmul:                                # @matmul
2      push    r15
3      push    r14
4      push    rbx
5      test    edi, edi
6      jle     .LBB0_7
7      mov     eax, edi
8      lea     r15, [4*rax]
9      xor     r8d, r8d
10     .LBB0_2:                            # =>This Loop Header: Depth=1
11         mov     r9, r8
12         imul    r9, rax
13         mov     r10, rdx
14         xor     r11d, r11d
15     .LBB0_3:                            # Parent Loop BB0_2 Depth=1
16         lea     r14, [r9 + r11]
17         vmovss  xmm0, dword ptr [rcx + 4*r14]    # xmm0 = mem[0],zero,zero,zero
18         mov     rbx, r10
19         xor     edi, edi
20     .LBB0_4:                            # Parent Loop BB0_2 Depth=1
21         vmovss  xmm1, dword ptr [rbx]            # xmm1 = mem[0],zero,zero,zero
22         vmulss  xmm1, xmm1, dword ptr [rsi + 4*rdi]
23         vaddss  xmm0, xmm0, xmm1
24         add     rdi, 1
25         add     rbx, r15
26         cmp     rax, rdi
27         jne     .LBB0_4
28         vmovss  dword ptr [rcx + 4*r14], xmm0
29         add     r11, 1
30         add     r10, 4
31         cmp     r11, rax
32         jne     .LBB0_3
33         add     r8, 1
34         add     rsi, r15
35         cmp     r8, rax
36         jne     .LBB0_2
37     .LBB0_7:
38         pop     rbx
39         pop     r14
40         pop     r15
41         ret

```

```

2      push    rbp
3      push    r15
4      push    r14
5      push    r13
6      push    r12
7      push    rbx
8      mov     quword ptr [rsp - 24], rcx    # 8-byte Spill
9      mov     quword ptr [rsp - 32], rdx    # 8-byte Spill
10     test    edi, edi
11     jle     .LBB0_12
12     mov     r9d, edi
13     mov     r8d, r9d
14     and     r8d, 8
15     lea     r13, [4*r9]
16     mov     r10, r9
17     shl     r10, 5
18     xor     eax, eax
19     vcvtsi2ss  xmm0, xmm0, xmm0
20     jmp     .LBB0_2
21     .LBB0_11:                            # In Loop: Header=BB0_2 Depth=1
22         mov     rcx, quword ptr [rsp - 16]    # 8-byte Reload
23         add     rcx, 1
24         add     r13, r13
25         cmp     rax, r9
26         je     .LBB0_12
27     .LBB0_2:                            # =>This Loop Header: Depth=1
28         mov     quword ptr [rsp - 16], rcx    # 8-byte Spill
29         imul    rcx, r9
30         mov     rcx, quword ptr [rsp - 24]    # 8-byte Reload
31         lea     rdx, [rcx + 4*rax]
32         mov     r15, quword ptr [rsp - 32]    # 8-byte Reload
33         xor     r12d, r12d
34         mov     quword ptr [rsp - 8], rdx    # 8-byte Spill
35         jmp     .LBB0_3
36     .LBB0_10:                            # In Loop: Header=BB0_3 Depth=2
37         vmovss  dword ptr [rdx + 4*r12], xmm1
38         add     r12, 1
39         add     r15, 4
40         cmp     r12, r9
41         je     .LBB0_11
42     .LBB0_3:                            # Parent Loop BB0_2 Depth=1
43         vmovss  xmm1, dword ptr [rdx + 4*r12]    # xmm1 = mem[0],zero,zero,zero
44         cmp     edi, 8
45         jne     .LBB0_5
46         xor     eax, eax
47         jmp     .LBB0_8
48     .LBB0_5:                            # In Loop: Header=BB0_3 Depth=2
49         vblendps  xmm1, xmm0, xmm1, 1    # xmm1 = xmm1[0],xmm0[1,2,3]
50         mov     rbx, r15
51         xor     eax, eax
52     .LBB0_6:                            # Parent Loop BB0_2 Depth=1
53         lea     rbp, [rbx + r13]
54         mov     r15, rbp
55         add     r13, r13
56         lea     rcx, [r11 + r13]
57         lea     rdx, [rcx + r13]
58         lea     r14, [rdx + r13]
59         vmovss  xmm2, dword ptr [r13 + rcx]    # xmm2 = mem[0],zero,zero,zero
60         vinsertps  xmm2, xmm2, dword ptr [r13 + rdx], 16 # xmm2 = xmm2[0],mem[0],xmm2[2,3]
61         vinsertps  xmm2, xmm2, dword ptr [r13 + r14], 32 # xmm2 = xmm2[0,1],mem[0],xmm2[3]
62         add     r14, r13
63         vinsertps  xmm2, xmm2, dword ptr [r13 + r14], 48 # xmm2 = xmm2[0,1,2],mem[0]
64         vmovss  xmm3, dword ptr [rbx]            # xmm3 = mem[0],zero,zero,zero
65         vinsertps  xmm3, xmm3, dword ptr [rbx + r13], 16 # xmm3 = xmm3[0],mem[0],xmm3[2,3]
66         vinsertps  xmm3, xmm3, dword ptr [r13 + rbp], 32 # xmm3 = xmm3[0,1],mem[0],xmm3[3]
67         vinsertps  xmm3, xmm3, dword ptr [r13 + r11], 48 # xmm3 = xmm3[0,1,2],mem[0]
68         vinsertf128  ymm0, ymm0, xmm2, 1
69         vmulps  ymm0, ymm0, ymmword ptr [rsi + 4*rax]
70         vaddps  ymm1, ymm1, ymm0
71         add     rax, 8
72         add     rbx, r10
73         cmp     r8, rax
74         jne     .LBB0_6
75         vextractf128  xmm2, ymm1, 1
76         vaddps  xmm1, xmm1, xmm2
77         vpermilpd  xmm2, xmm1, 1    # xmm2 = xmm1[1,0]
78         vaddps  xmm1, xmm1, xmm2
79         vmovshdup  xmm2, xmm1    # xmm2 = xmm1[1,1,3,3]
80         vaddss  xmm1, xmm1, xmm2
81         mov     rax, r8
82         cmp     r8, r9
83         mov     rdx, quword ptr [rsp - 8]    # 8-byte Reload
84         je     .LBB0_10
85     .LBB0_8:                            # In Loop: Header=BB0_3 Depth=2
86         mov     rcx, r9
87         imul    rcx, rax
88         lea     rbx, [r15 + 4*rcx]
89     .LBB0_9:                            # Parent Loop BB0_2 Depth=1
90         vmovss  xmm2, dword ptr [rbx]            # xmm2 = mem[0],zero,zero,zero
91         vmulss  xmm2, xmm2, dword ptr [rsi + 4*rax]
92         vaddss  xmm1, xmm1, xmm2
93         add     rax, 1
94         add     rbx, r13
95         cmp     r9, rax
96         jne     .LBB0_9
97         jmp     .LBB0_10
98     .LBB0_12:
99         pop     rbx
100        pop     r12
101        pop     r13
102        pop     r14
103        pop     r15
104        pop     rbp

```

8. Vectorization of Function Calls

- Works for a few built-in functions
- Note, need to add the exact library to the command line

```
#include <math.h>

void bar(float * restrict A, float * restrict B) {
    for (int i = 0; i < 1024; ++i) {
        B[i] = sinf(A[i]);
    }
}
```

8. Vectorization of Function Calls

```
bar:                                     # @bar
    pushq    %r15
    pushq    %r14
    pushq    %rbx
    movq     %rsi, %rbx
    movq     %rdi, %r14
    xorl     %r15d, %r15d
.LBB0_1:                                # =>This Inner Loop Header: Depth=1
    vmovups  (%r14,%r15,4), %ymm0
    callq    _ZGVdN8v_sinf@PLT
    vmovups  %ymm0, (%rbx,%r15,4)
    addq     $8, %r15
    cmpq     $1024, %r15                # imm = 0x400
    jne      .LBB0_1
    popq     %rbx
    popq     %r14
    popq     %r15
    vzeroupper
    retq
```

9. Different Hardware Versions

- Different Intel chips have different vector hardware
 - MMX, SSE, SSE2, SSE3, AVX, AVX2, AVX512, AVX512VNNI etc.

```
void bar(float * restrict A, float * restrict B) {  
    for (int i = 0; i < 4; ++i)  
        A[i] += B[i];  
}
```

9. Different Hardware Versions

The image displays four screenshots of the x86-64 clang 16.0.0 compiler output, arranged in a 2x2 grid. Each screenshot shows the assembly code for a function named 'bar' with a loop header 'Depth=1'. The instructions are color-coded: green for the loop header, blue for the loop body, and red for the loop footer. A red box highlights a specific instruction in each screenshot.

Top-Left Screenshot: Compiler flags: `-O3 -fno-vectorize -ffast-math -fno-unroll-loops`. The highlighted instruction is `incq %rax` (line 7).

Top-Right Screenshot: Compiler flags: `-O3 -msse2 -ffast-math -fno-unroll-loops`. The highlighted instruction is `addq $4, %rax` (line 8).

Bottom-Left Screenshot: Compiler flags: `-O3 -mavx2 -ffast-math -fno-unroll-loops`. The highlighted instruction is `addq $8, %rax` (line 7).

Bottom-Right Screenshot: Compiler flags: `-O3 -mavx512vnni -ffast-math -fno-unroll-loops`. The highlighted instruction is `addq $16, %rax` (line 7).

Aside: Do I Have SIMD Capabilities?

```
$ less /proc/cpuinfo
```

```
flags              : fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat p
se36 clflush dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx pdpe1gb rdtscp lm con
stant_tsc art arch_perfmon pebs bts rep_good nopl xtopology nonstop_tsc cpuid aperfmp
erf tsc_known_freq pni pclmulqdq dtes64 monitor ds_cpl vmx est tm2 ssse3 sdbg fma cx1
6 xtpr pdcm pcid sse4_1 sse4_2 x2apic movbe popcnt tsc_deadline_timer aes xsave avx f
16c rdrand lahf_lm abm 3dnowprefetch cpuid_fault epb invpcid_single pti ssbd ibrs ibp
b stibp tpr_shadow vnmi flexpriority ept vpid fsgsbase tsc_adjust bmi1 avx2 smep bmi2
erms invpcid mpx rdseed adx smap clflushopt intel_pt xsaveopt xsavec xgetbv1 xsaves
dtherm ida arat pln pts hwp hwp_notify hwp_act_window hwp_epp flush_l1d
```

Outline

- What is SIMD
- X86 SIMD Instructions
- Auto-vectorization (by compiler)
- **Hand Vectorization using Intrinsics: Blood Sweat and Tears**

Example: SAXPY

```
void saxpy(int n, float a, float *restrict x, float *restrict y) {  
    for (int i = 0; i < n; i++) {  
        y[i] = a * x[i] + y[i];  
    }  
}
```

```
void saxpy(int n, float a, float *x, float *y) {  
    int vec_len = 8;  
    for (int i = 0; i < n % vec_len; i++)  
        y[i] += a * x[i];  
  
    __m256 va = _mm256_set1_ps(a);  
    for (int i = n % vec_len; i < n; i += vec_len) {  
        __m256 vx = _mm256_loadu_ps(x + i);  
        __m256 vy = _mm256_loadu_ps(y + i);  
        __m256 ax = _mm256_mul_ps(va, vx);  
        __m256 axpy = _mm256_add_ps(ax, vy);  
        _mm256_storeu_ps(y + i, axpy);  
    }  
}
```

Example: SAXPY

```
void saxpy(int n, float a, float *restrict x, float *restrict y) {  
    for (int i = 0; i < n; i++) {  
        y[i] = a * x[i] + y[i];  
    }  
}
```

- Compiler can do this automatically too.
- Running time (n=1024)
 - Vectorized (256-bit vector) = **494** cycles
 - Scalar = **3309** cycles
 - ... about **6.7x** faster.

Vectorizing Reductions

```
float sum(int n, float *x) {  
    for (int i = 0; i < n; i++)  
        s += x[i];  
    return s;  
}
```

```
float sum(int n, float *x) {  
    int vec_len = 8;  
    float s = 0.0;  
    for (int i = 0; i < n % vec_len; i++)  
        s += x[i];  
  
    __m256 v = _mm256_set_ps(0., 0., 0., 0., 0., 0., 0., s);  
    for (int i = n % vec_len; i < n; i += vec_len) {  
        __m256 vx = _mm256_loadu_ps(x + i);  
        v = _mm256_add_ps(v, vx);  
    }  
    return reduce_vector(v);  
}
```

Vectorizing Reductions (cont.)

```
float reduce_vector(__m256 v) {  
    // v_0_4 = {v3, v2, v1, v0}  
    __m128 v_0_4 = _mm256_extracti128_si256(v, 0);  
    // v_4_8 = {v7, v6, v5, v4}  
    __m128 v_4_8 = _mm256_extracti128_si256(v, 1);  
    // t = {v3+v7, v2+v6, v1+v5, v0+v4}  
    __m128 t = _mm_add_ps(v_0_4, v_4_8);  
    // t1 = {v2+v6, v3+v7, v0+v4, v1+v5}  
    __m128 t1 = _mm_shuffle_ps(t, t, _MM_SHUFFLE(2, 3, 0, 1));  
    // t2 = {_, v2+v6+v3+v7, _, v0+v4+v1+v5}  
    __m128 t2 = _mm_add_ps(t, t1);  
    // t3 = {_, _, _, v2+v6+v3+v7}  
    __m128 t3 = _mm_shuffle_ps(t2, t2, _MM_SHUFFLE(0, 0, 0, 2));  
    // t4 = {_, _, _, v0+v4+v1+v5+v2+v6+v3+v7}  
    __m128 t4 = _mm_add_ps(t2, t3);  
    // Extract the first lane  
    return _mm_cvtss_f32(t4);  
}
```

Vectorizing Control Flow


```
void foo(float *restrict x, float *restrict y) {  
    for (int i = 0; i < 8; i++) {  
        if (x[i] < y[i])  
            x[i] += y[i];  
        else  
            x[i] -= y[i];  
    }  
}
```

```
__m256 x_vec    = _mm256_loadu_ps(x);  
__m256 y_vec    = _mm256_loadu_ps(y);  
__m256 x_lt_y   = _mm256_cmp_ps(x_vec, y_vec, _CMP_LT_OS);  
__m256 if_true  = _mm256_add_ps(x_vec, y_vec);  
__m256 if_false = _mm256_sub_ps(x_vec, y_vec);  
__m256 result   = _mm256_blendv_ps(if_true, if_false, x_lt_y);  
_mm256_storeu_ps(x, result);
```

Vectorizing the Dot Product

```
void dot_16x1x16(uint8_t *data, int8_t kernel[][4], int32_t *acc) {  
    for (int i = 0; i < 16; i++)  
        for (int j = 0; j < 4; j++)  
            acc[i] += data[j] * kernel[i][j];  
}
```

- Compiler vectorizes the inner loop
 - Similar to a reduction




2.2x faster than scalar

Vectorizing the Dot Product

```
void dot_16x1x16(uint8_t *data, int8_t kernel[][4], int32_t *acc) {  
    for (int i = 0; i < 16; i++)  
        for (int j = 0; j < 4; j++)  
            acc[i] += data[j] * kernel[i][j];  
}
```

- Compiler vectorizes the inner loop
 - Similar to a reduction




```
dot:                                # @dot  
    vpmovzxbd    (%rdi), %xmm0      # xmm0 = mem[0],zero,zero,zero,mem[1],zero,zero,zero,mem[2],zero,zero,zero,mem[3],zero,zero,zero  
    xorl        %eax, %eax  
    .LBB0_1:                          # =>This Inner Loop Header: Depth=1  
    vmovd       (%rdx,%rax,4), %xmm1 # xmm1 = mem[0],zero,zero,zero  
    vpmovsxbd    (%rsi,%rax,4), %xmm2  
    vpmaddwd     %xmm0, %xmm2, %xmm2  
    vpaddq      %xmm1, %xmm2, %xmm1  
    vpsltd      $238, %xmm1, %xmm2   # xmm2 = xmm1[2,3,2,3]  
    vpaddq      %xmm2, %xmm1, %xmm1  
    vpsltd      $85, %xmm1, %xmm2    # xmm2 = xmm1[1,1,1,1]  
    vpaddq      %xmm2, %xmm1, %xmm1  
    vmovd       %xmm1, (%rdx,%rax,4)  
    incq        %rax  
    cmpq        $16, %rax  
    jne         .LBB0_1  
    retq
```

< faster than scalar

Vectorizing the Dot Product

```
void dot_16x1x16(uint8_t *data, int8_t kernel[][4], int32_t *acc) {  
    for (int i = 0; i < 16; i++)  
        for (int j = 0; j < 4; j++)  
            acc[i] += data[j] * kernel[i][j];  
}
```

- Compiler vectorizes the inner loop
 - Similar to a reduction
- **Can we do better?**





2.2x faster than scalar

```
__m512i acc_vec = _mm512_loadu_epi32(acc);  
// {data0, data1, data2, data3, ..., data0, data1, data2, data3}  
__m512i data_vec = _mm512_set1_epi32(*(int32_t)data);  
__m512i kernel_vec = _mm512_loadu_epi32(&kernel);  
_mm512_dpbusd_epi32(acc_vec, data_vec, kernel_vec);
```

Vectorizing the Dot Product

```
void dot_16x1x16(uint8_t *data, int8_t kernel[][4], int32_t *acc) {
    for (int i = 0; i < 16; i++)
        for (int j = 0; j < 4; j++)
            acc[i] += data[j] * kernel[i][j];
}
```

VPDPBUSD



2.2x faster than scalar

```
__m512i acc_vec = _mm512_loadu_epi32(acc);
// {data0, data1, data2, data3, ..., data0, data1, data2, data3}
__m512i data_vec = _mm512_set1_epi32(*(int32_t *)data);
__m512i kernel_vec = _mm512_loadu_epi32(&kernel);
_mm512_dpbasd_epi32(acc_vec, data_vec, kernel_vec);
```

11x faster than scalar

Summary

- **Vector** hardware is great for performance (**data parallelism**)
 - SIMD: Single Instruction Multiple Data
 - Architecture is changing / improving quickly (latest SIMD generation)
- **Compilers** can vectorize your code
 - However, compilers are finicky
 - May have to coax the compiler
- Can also use **intrinsics** to hand vectorize
 - ‘Assembly programming’
 - Hard to keep-up with hardware updates