



LECTURE 6

Multicore Programming


Xuhao Chen

February 7, 2025


From Single-Core to Multi-Core

- Why Multicore?


Single-core Processor



Multi-core Processor



Moore's Law



Processor data from Stanford's CPU DB [DKM12].


Power Wall at 4 GHz

- Power $\approx$ Voltage² * Frequency
- The end of Dennard Scaling



Vendor Solution: Multicore


- To scale performance, vendors put many processing cores on the chip
- Each generation of Moore's Law potentially **doubles the number of cores**




Intel Core i7 3960X
(Sandy Bridge E), 2011

- 6 cores
- 3.3 GHz
- 15-MB L3 cache

Historical Software Scaling



Ideal Multicore Scaling



Unfortunately, not so simple...



Free Lunch is Over *




* Herb Sutter, The Free Lunch Is Over: A Fundamental Turn Toward Concurrency in Software, 2010

Serial vs. Parallel Programming


Parallel programming is hard!




Load Imbalance



Synchronization & Communication



- Synchronization



- Communication


OUTLINE

- Shared-Memory Hardware
- Concurrency Platforms
 - ▣ Pthreads (and WinAPI Threads)
 - ▣ Threading Building Blocks (TBB)
 - ▣ OpenMP
 - ▣ Cilk

SHARED-MEMORY HARDWARE




Abstract Multicore Architecture




Shared Memory vs. Distributed Memory


Shared Memory




Distributed Memory




Abstract Multicore Architecture




Cache Coherence




Cache Coherence




Cache Coherence




Cache Coherence



Cache Coherence




Cache Coherence



MSI Protocol

Each cache line is labeled with a state:

- **M**: cache block has been **modified**. No other caches contain this block in **M** or **S**.
- **S**: other caches may be **sharing** this block.
- **I**: cache block is **invalid** (same as not there).




Before a cache modifies a location, the hardware first invalidates all other copies.

MSI Protocol

Each cache line is labeled with a state:


- **M**: cache block has been **modified**. No other caches contain this block in **M** or **S**.
- **S**: other caches may be **sharing** this block.
- **I**: cache block is **invalid** (same as not there).



MSI Protocol

Each cache line is labeled with a state:


- **M**: cache block has been **modified**. No other caches contain this block in **M** or **S**.
- **S**: other caches may be **sharing** this block.
- **I**: cache block is **invalid** (same as not there).



MSI Protocol

Each cache line is labeled with a state:


- **M**: cache block has been **modified**. No other caches contain this block in **M** or **S**.
- **S**: other caches may be **sharing** this block.
- **I**: cache block is **invalid** (same as not there).



MSI Protocol

Each cache line is labeled with a state:


- **M**: cache block has been **modified**. No other caches contain this block in **M** or **S**.
- **S**: other caches may be **sharing** this block.
- **I**: cache block is **invalid** (same as not there).



MSI Protocol

Each cache line is labeled with a state:

- **M**: cache block has been **modified**. No other caches contain this block in **M** or **S**.
- **S**: other caches may be **sharing** this block.
- **I**: cache block is **invalid** (same as not there).



CONCURRENCY PLATFORMS



Concurrency Platforms

- Programming directly on processor cores is **painful** and **error-prone**
- A **concurrency platform** helps
 - abstract processor cores
 - handle synchronization and communication
 - perform load balancing
- Examples
 - Pthreads and WinAPI threads
 - Threading Building Blocks (TBB)
 - OpenMP
 - **Cilk**

Fibonacci Numbers

The **Fibonacci numbers** are the sequence $\langle 0, 1, 1, 2, 3, 5, 8, 13, 21, 34, \dots \rangle$, where each number is the sum of the previous two.

Recurrence:

$$F_0 = 0,$$

$$F_1 = 1,$$

$$F_n = F_{n-1} + F_{n-2} \text{ for } n > 1.$$



The sequence is named after **Leonardo di Pisa** (1170–1250 A.D.), also known as Fibonacci, a contraction of filius Bonaccii —“son of Bonaccio.” Fibonacci’s 1202 book *Liber Abaci* introduced the sequence to Western mathematics, although it had previously been discovered by Indian mathematicians.

Fibonacci Program

```
#include <inttypes.h>
#include <stdio.h>
#include <stdlib.h>


int64_t fib(int64_t n) {
    if (n < 2) {
        return n;
    } else {
        int64_t x = fib(n-1);
        int64_t y = fib(n-2);
        return (x + y);
    }
}

int main(int argc, char *argv[]) {
    int64_t n = atoi(argv[1]);
    int64_t result = fib(n);
    printf("Fibonacci of %" PRIu64 " is %" PRIu64 ".\n",
           n, result);
    return 0;
}
```

Disclaimer to Algorithms Police

This recursive program is a poor way to compute the n th Fibonacci number, but it provides a good didactic example.

Fibonacci Execution



Key idea for parallelization
`fib(n-1)` and `fib(n-2)` can
be executed simultaneously
without mutual interference

```
int64_t fib(int64_t n) {  
    if (n < 2) {  
        return n;  
    } else {  
        int64_t x = fib(n-1);  
        int64_t y = fib(n-2);  
        return (x + y);  
    }  
}
```

CONCURRENCY PLATFORMS

- **Pthreads (& WinAPI Threads)**
- Threading Building Blocks
- OpenMP
- Cilk



Pthreads*

- Standard API for threading specified by ANSI/IEEE POSIX 1003.1-2008
- Do-it-yourself concurrency platform
- Built as a library of functions with “special” non-C semantics
- Each thread implements an abstraction of a processor, which are multiplexed onto machine resources
- Threads communicate through shared memory
- Library functions mask the protocols involved in inter-thread coordination



*WinAPI threads provide similar functionality.

Key Pthread Functions

```
int pthread_create(  
    pthread_t *thread,  
        //returned identifier for the new thread  
    const pthread_attr_t *attr,  
        //object to set thread attributes (NULL for default)  
    void *(*func)(void *),  
        //routine executed after creation  
    void *arg  
        //a single argument passed to func  
) //returns error status
```

```
int pthread_join(  
    pthread_t thread,  
        //identifier of thread to wait for  
    void **status  
        //terminating thread's status (NULL to ignore)  
) //returns error status
```

fork-join Programming Model



Pthread fib Implementation

```
#include <inttypes.h>
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

int64_t fib(int64_t n) {
    if (n < 2) {
        return n;
    } else {
        int64_t x = fib(n-1);
        int64_t y = fib(n-2);
        return (x + y);
    }
}

typedef struct {
    int64_t input;
    int64_t output;
} thread_args;

void *thread_func(void *ptr) {
    int64_t i = ((thread_args *) ptr)->input;
    ((thread_args *) ptr)->output = fib(i);
    return NULL;
}
```

```
int main(int argc, char *argv[]) {
    pthread_t thread;
    thread_args args;
    int status;
    int64_t result;

    if (argc < 2) { return 1; }
    int64_t n = strtoul(argv[1], NULL, 0);
    if (n < 30) {
        result = fib(n);
    } else {
        args.input = n-1;
        status = pthread_create(&thread,
                                NULL,
                                thread_func,
                                (void*) &args);

        // main can continue executing
        if (status != NULL) { return 1; }
        result = fib(n-2);
        // wait for the thread to terminate
        status = pthread_join(thread, NULL);
        if (status != NULL) { return 1; }
        result += args.output;
    }
    printf("Fibonacci of %" PRId64 " is %" PRId64 ".\n",
           n, result);
    return 0;
}
```

Pthread fib Implementation

```
#include <inttypes.h>
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

int64_t fib(int64_t n) {
    if (n < 2) {
        return n;
    } else {
        int64_t x = fib(n-1);
        int64_t y = fib(n-2);
        return (x + y);
    }
}

typedef struct {
    int64_t input;
    int64_t output;
} thread_args;

void *thread_func(void *ptr) {
    int64_t i = ((thread_args *) ptr)->input;
    ((thread_args *) ptr)->output = fib(i);
    return NULL;
}
```

Original code

```
int main(int argc, char *argv[]) {
    pthread_t thread;
    thread_args args;
    int status;
    int64_t result;

    if (argc < 2) { return 1; }
    int64_t n = strtoul(argv[1], NULL, 0);
    if (n < 30) {
        result = fib(n);
    } else {
        args.input = n-1;
        status = pthread_create(&thread,
                                NULL,
                                thread_func,
                                (void*) &args);

        // main can continue executing
        if (status != NULL) { return 1; }
        result = fib(n-2);
        // wait for the thread to terminate
        status = pthread_join(thread, NULL);
        if (status != NULL) { return 1; }
        result += args.output;
    }
    printf("Fibonacci of %" PRId64 " is %" PRId64 ".\n",
           n, result);
    return 0;
}
```

Pthread fib Implementation

```
#include <inttypes.h>
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

int64_t fib(int64_t n) {
    if (n < 2) {
        return n;
    } else {
        int64_t x = fib(n-1);
        int64_t y = fib(n-2);
        return (x + y);
    }
}

typedef struct {
    int64_t input;
    int64_t output;
} thread_args;

void *thread_func(void *ptr) {
    int64_t i = ((thread_args *) ptr)->input;
    ((thread_args *) ptr)->output = fib(i);
    return NULL;
}
```

Structure for
thread arguments

```
int main(int argc, char *argv[]) {
    pthread_t thread;
    thread_args args;
    int status;
    int64_t result;

    if (argc < 2) { return 1; }
    int64_t n = strtoul(argv[1], NULL, 0);
    if (n < 30) {
        result = fib(n);
    } else {
        args.input = n-1;
        status = pthread_create(&thread,
                                NULL,
                                thread_func,
                                (void*) &args);

        // main can continue executing
        if (status != NULL) { return 1; }
        result = fib(n-2);
        // wait for the thread to terminate
        status = pthread_join(thread, NULL);
        if (status != NULL) { return 1; }
        result += args.output;
    }
    printf("Fibonacci of %" PRId64 " is %" PRId64 ".\n",
          n, result);
    return 0;
}
```


Pthread **fib** Implementation

```
#include <inttypes.h>
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

int64_t fib(int64_t n) {
    if (n < 2) {
        return n;
    } else {
        int64_t x = fib(n-1);
        int64_t y = fib(n-2);
        return (x + y);
    }
}

typedef struct {
    int64_t input;
    int64_t output;
} thread_args;

void *thread_func(void *ptr) {
    int64_t i = ((thread_args *) ptr)->input;
    ((thread_args *) ptr)->output = fib(i);
    return NULL;
}
```

Function
called when
thread is
created

```
int main(int argc, char *argv[]) {
    pthread_t thread;
    thread_args args;
    int status;
    int64_t result;

    if (argc < 2) { return 1; }
    int64_t n = strtoul(argv[1], NULL, 0);
    if (n < 30) {
        result = fib(n);
    } else {
        args.input = n-1;
        status = pthread_create(&thread,
                                NULL,
                                thread_func,
                                (void*) &args);

        // main can continue executing
        if (status != NULL) { return 1; }
        result = fib(n-2);
        // wait for the thread to terminate
        status = pthread_join(thread, NULL);
        if (status != NULL) { return 1; }
        result += args.output;
    }
    printf("Fibonacci of %" PRId64 " is %" PRId64 ".\n",
          n, result);
    return 0;
}
```

Pthread **fib** Implementation

```
#include <inttypes.h>
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

int64_t fib(int64_t n) {
    if (n < 2) {
        return n;
    } else {
        int64_t x = fib(n-1);
        int64_t y = fib(n-2);
        return (x + y);
    }
}

typedef struct {
    int64_t input;
    int64_t output;
} thread_args;

void *thread_func(void *ptr) {
    int64_t i = ((thread_args *) ptr)->input;
    ((thread_args *) ptr)->output = fib(i);
    return NULL;
}
```

No point in creating thread if there isn't enough to do

```
int main(int argc, char *argv[]) {
    pthread_t thread;
    thread_args args;
    int status;
    int64_t result;

    if (argc < 2) { return 1; }
    int64_t n = strtoul(argv[1], NULL, 0);
    if (n < 30) {
        result = fib(n);
    } else {
        args.input = n-1;
        status = pthread_create(&thread,
                                NULL,
                                thread_func,
                                (void*) &args);

        // main can continue executing
        if (status != NULL) { return 1; }
        result = fib(n-2);
        // wait for the thread to terminate
        status = pthread_join(thread, NULL);
        if (status != NULL) { return 1; }
        result += args.output;
    }
    printf("Fibonacci of %" PRId64 " is %" PRId64 ".\n",
          n, result);
    return 0;
}
```

Pthread fib Implementation

```
#include <inttypes.h>
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

int64_t fib(int64_t n) {
    if (n < 2) {
        return n;
    } else {
        int64_t x = fib(n-1);
        int64_t y = fib(n-2);
        return (x + y);
    }
}

typedef struct {
    int64_t input;
    int64_t output;
} thread_args;

void *thread_func(void *ptr) {
    int64_t i = ((thread_args *) ptr)->input;
    ((thread_args *) ptr)->output = fib(i);
    return NULL;
}
```

```
int main(int argc, char *argv[]) {
    pthread_t thread;
    thread_args args;
    int status;
    int64_t result;

    if (argc < 2) { return 1; }
    int64_t n = strtoul(argv[1], NULL, 10);
    if (n < 30) {
        result = fib(n);
    } else {
        args.input = n-1;
        status = pthread_create(&thread,
                                NULL,
                                thread_func,
                                (void*) &args);

        // main can continue executing
        if (status != NULL) { return 1; }
        result = fib(n-2);
        // wait for the thread to terminate
        status = pthread_join(thread, NULL);
        if (status != NULL) { return 1; }
        result += args.output;
    }
    printf("Fibonacci of %" PRId64 " is %" PRId64 ".\n",
           n, result);
    return 0;
}
```

Marshal input
argument to
thread

Pthread fib Implementation

```
#include <inttypes.h>
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

int64_t fib(int64_t n) {
    if (n < 2) {
        return n;
    } else {
        int64_t x = fib(n-1);
        int64_t y = fib(n-2);
        return (x + y);
    }
}

typedef struct {
    int64_t input;
    int64_t output;
} thread_args;

void *thread_func(void *ptr) {
    int64_t i = ((thread_args *) ptr)->input;
    ((thread_args *) ptr)->output = fib(i);
    return NULL;
}
```

Create thread to
execute **fib(n-1)**

```
int main(int argc, char *argv[]) {
    pthread_t thread;
    thread_args args;
    int status;
    int64_t result;

    if (argc < 2) { return 1; }
    int64_t n = strtoul(argv[1], NULL, 0);
    if (n < 30) {
        result = fib(n);
    } else {
        args.input = n-1;
        status = pthread_create(&thread,
                                NULL,
                                thread_func,
                                (void*) &args);

        // main can continue executing
        if (status != NULL) { return 1; }
        result = fib(n-2);
        // wait for the thread to terminate
        status = pthread_join(thread, NULL);
        if (status != NULL) { return 1; }
        result += args.output;
    }
    printf("Fibonacci of %" PRId64 " is %" PRId64 ".\n",
          n, result);
    return 0;
}
```

Pthread fib Implementation

```
#include <inttypes.h>
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

int64_t fib(int64_t n) {
    if (n < 2) {
        return n;
    } else {
        int64_t x = fib(n-1);
        int64_t y = fib(n-2);
        return (x + y);
    }
}

typedef struct {
    int64_t input;
    int64_t output;
} thread_args;

void *thread_func(
    int64_t i = ((thread_args *) ptr)->input;
    ((thread_args *) ptr)->output = fib(i);
    return NULL;
}
```

Main program
runs **fib(n-2)**
in parallel.

```
int main(int argc, char *argv[]) {
    pthread_t thread;
    thread_args args;
    int status;
    int64_t result;

    if (argc < 2) { return 1; }
    int64_t n = strtoul(argv[1], NULL, 0);
    if (n < 30) {
        result = fib(n);
    } else {
        args.input = n-1;
        status = pthread_create(&thread,
                                NULL,
                                thread_func,
                                (void*) &args);

        // main can continue executing
        if (status != NULL) { return 1; }
        result = fib(n-2);
        // wait for the thread to terminate
        status = pthread_join(thread, NULL);
        if (status != NULL) { return 1; }
        result += args.output;
    }
    printf("Fibonacci of %" PRId64 " is %" PRId64 ".\n",
          n, result);
    return 0;
}
```


Pthread **fib** Implementation

```
#include <inttypes.h>
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

int64_t fib(int64_t n) {
    if (n < 2) {
        return n;
    } else {
        int64_t x = fib(n-1);
        int64_t y = fib(n-2);
        return (x + y);
    }
}

typedef struct {
    int64_t input;
    int64_t output;
} thread_args;

void *thread_func(void *ptr) {
    int64_t i = ((thread_args *) ptr)->input;
    ((thread_args *) ptr)->output = fib(i);
    return NULL;
}
```

Block until the
auxiliary thread
finishes.

```
int main(int argc, char *argv[]) {
    pthread_t thread;
    thread_args args;
    int status;
    int64_t result;

    if (argc < 2) { return 1; }
    int64_t n = strtoul(argv[1], NULL, 0);
    if (n < 30) {
        result = fib(n);
    } else {
        args.input = n-1;
        status = pthread_create(&thread,
                                NULL,
                                thread_func,
                                (void*) &args);

        // main can continue executing
        if (status != NULL) { return 1; }
        result = fib(n-2);
        // wait for the thread to terminate
        status = pthread_join(thread, NULL);
        if (status != NULL) { return 1; }
        result += args.output;
    }
    printf("Fibonacci of %" PRId64 " is %" PRId64 ".\n",
          n, result);
    return 0;
}
```

Pthread fib Implementation

```
#include <inttypes.h>
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

int64_t fib(int64_t n) {
    if (n < 2) {
        return n;
    } else {
        int64_t x = fib(n-1);
        int64_t y = fib(n-2);
        return (x + y);
    }
}

typedef struct {
    int64_t input;
    int64_t output;
} thread_args;

void *thread_func(void *ptr) {
    int64_t i = ((thread_args *) ptr)->input;
    ((thread_args *) ptr)->output = fib(i);
    return NULL;
}
```

Add the results together to produce the final output.

```
int main(int argc, char *argv[]) {
    pthread_t thread;
    thread_args args;
    int status;
    int64_t result;

    if (argc < 2) { return 1; }
    int64_t n = strtoul(argv[1], NULL, 0);
    if (n < 30) {
        result = fib(n);
    } else {
        args.input = n-1;
        status = pthread_create(&thread,
                                NULL,
                                thread_func,
                                (void*) &args);

        // main can continue executing
        if (status != NULL) { return 1; }
        result = fib(n-2);
        // wait for the thread to terminate
        status = pthread_join(thread, NULL);
        if (status != NULL) { return 1; }
        result += args.output;
    }
    printf("Fibonacci of %" PRId64 " is %" PRId64 ".\n",
           n, result);
    return 0;
}
```

Issues with Pthreads

Overhead	The cost of creating a thread $>10^4$ cycles $\Rightarrow$ coarse-grained concurrency. (Thread pools can help.)

Issues with Pthreads

Overhead	The cost of creating a thread $>10^4$ cycles $\Rightarrow$ coarse-grained concurrency. (Thread pools can help.)
Scalability	Fibonacci code gets at most about 1.5 speedup for 2 cores. Need a rewrite for more cores

Issues with Pthreads

Overhead	The cost of creating a thread $>10^4$ cycles $\Rightarrow$ coarse-grained concurrency. (Thread pools can help.)
Scalability	Fibonacci code gets at most about 1.5 speedup for 2 cores. Need a rewrite for more cores
Modularity	The Fibonacci logic is no longer neatly encapsulated in the <code>fib()</code> function

Issues with Pthreads

Overhead	The cost of creating a thread $>10^4$ cycles $\Rightarrow$ coarse-grained concurrency. (Thread pools can help.)
Scalability	Fibonacci code gets at most about 1.5 speedup for 2 cores. Need a rewrite for more cores
Modularity	The Fibonacci logic is no longer neatly encapsulated in the <code>fib()</code> function
Code Simplicity	Programmers must marshal arguments (shades of 1957!) and engage in error-prone protocols for load-balance


CONCURRENCY PLATFORMS

- Pthreads (& WinAPI Threads)
- **Threading Building Blocks**
- OpenMP
- Cilk



Threading Building Blocks

- Developed by Intel, as a **C++ library** that runs on native threads
- Programmer specifies **tasks** rather than threads
- Tasks are automatically load balanced across the threads using a **work-stealing** algorithm inspired by Cilk at MIT
- Focus on performance



Fibonacci in TBB

```
using namespace tbb;
class FibTask: public task {
public:
    const int64_t n;
    int64_t* const sum;
    FibTask(int64_t n_, int64_t* sum_) :
        n(n_), sum(sum_) {}

    task* execute() {
        if( n < 2 ) {
            *sum = n;
        } else {
            int64_t x, y;
            FibTask& a = *new( allocate_child() )
                          FibTask(n-1, &x);
            FibTask& b = *new( allocate_child() )
                          FibTask(n-2, &y);

            set_ref_count(3);
            spawn(b);
            spawn_and_wait_for_all(a);
            *sum = x + y;
        }
        return NULL;
    }
};
```

```
#include <cstdint>
#include <iostream>
#include "tbb/task.h"

int main(int argc, char *argv[]) {
    int64_t res;
    if (argc < 2) { return 1; }
    int64_t n =
        strtoul(argv[1], NULL, 0);
    FibTask& a = *new(task::allocate_root())
                 FibTask(n, &res);
    task::spawn_root_and_wait(a);

    std::cout << "Fibonacci of " << n
               << " is " << res << std::endl;
    return 0;
}
```

Fibonacci in TBB

```
using namespace tbb;
class FibTask: public task {
public:
    const int64_t n;
    int64_t* const sum;
    FibTask(int64_t n_, int64_t* sum_) :
        n(n_), sum(sum_) {}

    task* execute() {
        if( n < 2 ) {
            *sum = n;
        } else {
            int64_t x, y;
            FibTask& a = *new( allocate_child() )
                          FibTask(n-1, &x);
            FibTask& b = *new( allocate_child() )
                          FibTask(n-2, &y);

            set_ref_count(3);
            spawn(b);
            spawn_and_wait_for_all(a);
            *sum = x + y;
        }
        return NULL;
    }
};
```

A computation
organized as
explicit tasks.

```
#include <cstdint>
#include <iostream>
#include "tbb/task.h"

int main(int argc, char *argv[]) {
    int64_t res;
    if (argc < 2) { return 1; }
    int64_t n =
        strtoul(argv[1], NULL, 0);
    FibTask& a = *new(task::allocate_root())
                 FibTask(n, &res);
    task::spawn_root_and_wait(a);

    std::cout << "Fibonacci of " << n
               << " is " << res << std::endl;
    return 0;
}
```

Fibonacci in TBB

```
using namespace tbb;
class FibTask: public task {
public:
    const int64_t n;
    int64_t* const sum;
    FibTask(int64_t n_, int64_t* sum_) :
        n(n_), sum(sum_) {}

    task* execute() {
        if( n < 2 ) {
            *sum = n;
        } else {
            int64_t x, y;
            FibTask& a = *new( allocate_child() )
                          FibTask(n-1, &x);
            FibTask& b = *new( allocate_child() )
                          FibTask(n-2, &y);

            set_ref_count(3);
            spawn(b);
            spawn_and_wait_for_all(a);
            *sum = x + y;
        }
        return NULL;
    }
};
```

FibTask has an
input parameter **n**
and an output
parameter **sum**.

```
#include <cstdint>
#include <iostream>
#include "tbb/task.h"

int main(int argc, char *argv[]) {
    int64_t res;
    if (argc < 2) { return 1; }
    int64_t n =
        strtoul(argv[1], NULL, 0);
    FibTask& a = *new(task::allocate_root())
                 FibTask(n, &res);
    task::spawn_root_and_wait(a);

    std::cout << "Fibonacci of " << n
               << " is " << res << std::endl;
    return 0;
}
```


Fibonacci in TBB

```
using namespace tbb;
class FibTask: public task {
public:
    const int64_t n;
    int64_t* const sum;
    FibTask(int64_t n_, int64_t* s) :
        n(n_), sum(s) {}

    task* execute() {
        if( n < 2 ) {
            *sum = n;
        } else {
            int64_t x, y;
            FibTask& a = *new( allocate_child() )
                          FibTask(n-1, &x);
            FibTask& b = *new( allocate_child() )
                          FibTask(n-2, &y);

            set_ref_count(3);
            spawn(b);
            spawn_and_wait_for_all(a);
            *sum = x + y;
        }
        return NULL;
    }
};
```

The `execute()` function performs the computation when the task is started.

```
#include <cstdint>
#include <iostream>
#include "tbb/task.h"

int main(int argc, char *argv[]) {
    int64_t res;
    if (argc < 2) { return 1; }
    int64_t n =
        strtoul(argv[1], NULL, 0);
    FibTask& a = *new(task::allocate_root())
                 FibTask(n, &res);
    task::spawn_root_and_wait(a);

    std::cout << "Fibonacci of " << n
               << " is " << res << std::endl;
    return 0;
}
```

Fibonacci in TBB

```
using namespace tbb;
class FibTask: public task {
public:
    const int64_t n;
    int64_t* const sum;
    FibTask(int64_t n_, int64_t* sum_) :
        n(n_), sum(sum_) {}

    task* execute() {
        if( n < 2 ) {
            *sum = n;
        } else {
            int64_t x, y;
            FibTask& a = *new( allocate_child() )
                          FibTask(n-1, &x);
            FibTask& b = *new( allocate_child() )
                          FibTask(n-2, &y);

            set_ref_count(3);
            spawn(b);
            spawn_and_wait_for_all(a);
            *sum = x + y;
        }
        return NULL;
    }
};
```

Recursively
create two
child tasks
a and **b**.

```
#include <cstdint>
#include <iostream>
#include "tbb/task.h"

int main(int argc, char *argv[]) {
    int64_t res;
    if (argc < 2) { return 1; }
    int64_t n =
        strtoul(argv[1], NULL, 0);
    FibTask& a = *new(task::allocate_root())
                 FibTask(n, &res);
    task::spawn_root_and_wait(a);

    std::cout << "Fibonacci of " << n
               << " is " << res << std::endl;
    return 0;
}
```

Fibonacci in TBB

```
using namespace tbb;
class FibTask: public task {
public:
    const int64_t n;
    int64_t* const sum;
    FibTask(int64_t n_, int64_t* sum_) :
        n(n_), sum(sum_) {}

    task* execute() {
        if( n < 2 ) {
            *sum = n;
        } else {
            int64_t x, y;
            FibTask& a = *new( allocate_child() )
                          FibTask(n-1, &x);
            FibTask& b = *new( allocate_child() )
                          FibTask(n-2, &y);

            set_ref_count(3);
            spawn(b);
            spawn_and_wait_for_all(a);
            *sum = x + y;
        }
        return NULL;
    }
};
```

Set the number of tasks to wait for (2 children + 1 implicit for bookkeeping).

```
#include <stdint>
#include <iostream>
#include "tbb/task.h"

int main(int argc, char *argv[]) {
    int64_t res;
    if (argc < 2) { return 1; }
    int64_t n =
        strtoul(argv[1], NULL, 0);
    FibTask& a = *new(task::allocate_root())
                 FibTask(n, &res);
    task::spawn_root_and_wait(a);

    std::cout << "Fibonacci of " << n
               << " is " << res << std::endl;
    return 0;
}
```

Fibonacci in TBB

```
using namespace tbb;
class FibTask: public task {
public:
    const int64_t n;
    int64_t* const sum;
    FibTask(int64_t n_, int64_t* sum_) :
        n(n_), sum(sum_) {}

    task* execute() {
        if( n < 2 ) {
            *sum = n;
        } else {
            int64_t x, y;
            FibTask& a = *new( allocate_child() )
                          FibTask(n-1, &x);
            FibTask& b = *new( allocate_child() )
                          FibTask(n-2, &y);

            set_ref_count(3);
            spawn(b);
            spawn_and_wait_for_all(a);
            *sum = x + y;
        }
        return NULL;
    }
};
```

Start task **b**.

```
#include <cstdint>
#include <iostream>
#include "tbb/task.h"

int main(int argc, char *argv[]) {
    int64_t res;
    if (argc < 2) { return 1; }
    int64_t n =
        strtoul(argv[1], NULL, 0);
    FibTask& a = *new(task::allocate_root())
                 FibTask(n, &res);
    task::spawn_root_and_wait(a);

    std::cout << "Fibonacci of " << n
               << " is " << res << std::endl;
    return 0;
}
```

Fibonacci in TBB

```
using namespace tbb;
class FibTask: public task {
public:
    const int64_t n;
    int64_t* const sum;
    FibTask(int64_t n_, int64_t* sum_) :
        n(n_), sum(sum_) {}

    task* execute() {
        if( n < 2 ) {
            *sum = n;
        } else {
            int64_t x, y;
            FibTask& a = *new( allocate_child() )
                          FibTask(n-1, &x);
            FibTask& b = *new( allocate_child() )
                          FibTask(n-2, &y);

            set_ref_count(3);
            spawn(b);
            spawn_and_wait_for_all(a);
            *sum = x + y;
        }
        return NULL;
    }
};
```

Start task **a**, and wait for both
a and **b** to finish.

```
#include <stdint>
#include <iostream>
#include "tbb/task.h"

int main(int argc, char *argv[]) {
    int64_t res;
    if (argc < 2) { return 1; }
    int64_t n =
        strtoul(argv[1], NULL, 0);
    FibTask& a = *new(task::allocate_root())
                 FibTask(n, &res);
    task::spawn_root_and_wait(a);

    std::cout << "Fibonacci of " << n
               << " is " << res << std::endl;
    return 0;
}
```

Fibonacci in TBB

```
using namespace tbb;
class FibTask: public task {
public:
    const int64_t n;
    int64_t* const sum;
    FibTask(int64_t n_, int64_t* sum_) :
        n(n_), sum(sum_) {}

    task* execute() {
        if( n < 2 ) {
            *sum = n;
        } else {
            int64_t x, y;
            FibTask& a = *new( allocate_child() )
                          FibTask(n-1, &x);
            FibTask& b = *new( allocate_child() )
                          FibTask(n-2, &y);

            set_ref_count(3);
            spawn(b);
            spawn_and_wait_for_all(a);
            *sum = x + y;
        }
        return NULL;
    }
};
```

Add the results together to produce the final output.

```
#include <stdint>
#include <iostream>
#include "tbb/task.h"

int main(int argc, char *argv[]) {
    int64_t res;
    if (argc < 2) { return 1; }
    int64_t n =
        strtoul(argv[1], NULL, 0);
    FibTask& a = *new(task::allocate_root())
                 FibTask(n, &res);
    task::spawn_root_and_wait(a);

    std::cout << "Fibonacci of " << n
               << " is " << res << std::endl;
    return 0;
}
```


Fibonacci in TBB

```
using namespace tbb;
class FibTask: public task {
public:
    const int64_t n;
    int64_t* const sum;
    FibTask(int64_t n_, int64_t* sum_) :
        n(n_), sum(sum_) {}
```

```
    task* execute() {
        if( n < 2 ) {
            *sum = n;
        } else {
            int64_t x, y;
            FibTask& a = *new( allocate_child() )
                          FibTask(n-1, &x);
            FibTask& b = *new( allocate_child() )
                          FibTask(n-2, &y);

            set_ref_count(3);
            spawn(b);
            spawn_and_wait_for_all(a);
            *sum = x + y;
        }
        return NULL;
    }
};
```

Create root task;
spawn and wait.

```
#include <cstdint>
#include <iostream>
#include "tbb/task.h"

int main(int argc, char *argv[]) {
    int64_t res;
    if (argc < 2) { return 1; }
    int64_t n =
        strtoul(argv[1], NULL, 0);
    FibTask& a = *new(task::allocate_root())
                 FibTask(n, &res);
    task::spawn_root_and_wait(a);

    std::cout << "Fibonacci of " << n
               << " is " << res << std::endl;
    return 0;
}
```

Other TBB Features

- TBB provides many C++ templates to express common patterns simply
 - ▣ e.g. `parallel_for` for loop parallelism,
 - ▣ e.g. `parallel_reduce` for data aggregation,
 - ▣ e.g. `pipeline` and `filter` for software pipelining
- TBB provides `concurrent container` classes, which allow multiple threads to safely access and update items in the container concurrently
- TBB also provides a variety of `mutual-exclusion` library functions, including `locks` and `atomic updates`

CONCURRENCY PLATFORMS

- Pthreads (& WinAPI Threads)
- Threading Building Blocks
- **OpenMP**
- Cilk



OpenMP

- Specification by an industry consortium
- Several compilers available, both open-source and proprietary, including GCC, ICC, Clang, and Visual Studio
- Linguistic extensions to C/C++ and Fortran in the form of compiler **pragmas**
- Runs on top of native threads
- Supports loop parallelism, task parallelism, and pipeline parallelism



Fibonacci in OpenMP

```
int64_t fib(int64_t n) {  
    if (n < 2) {  
        return n;  
    } else {  
        int64_t x, y;  
        #pragma omp task shared(x,n)  
        x = fib(n-1);  
        #pragma omp task shared(y,n)  
        y = fib(n-2);  
        #pragma omp taskwait  
        return (x + y);  
    }  
}
```

Compiler directive.

Fibonacci in OpenMP

```
int64_t fib(int64_t n) {  
    if (n < 2) {  
        return n;  
    } else {  
        int64_t x, y;  
        #pragma omp task shared(x,n)  
        x = fib(n-1);  
        #pragma omp task shared(y,n)  
        y = fib(n-2);  
        #pragma omp taskwait  
        return (x + y);  
    }  
}
```

The following statement is an independent task.

Fibonacci in OpenMP

```
int64_t fib(int64_t n) {  
    if (n < 2) {  
        return n;  
    } else {  
        int64_t x, y;  
        #pragma omp task shared(x,n)  
        x = fib(n-1);  
        #pragma omp task shared(y,n)  
        y = fib(n-2);  
        #pragma omp taskwait  
        return (x + y);  
    }  
}
```

Sharing of memory is
managed explicitly.

Fibonacci in OpenMP

```
int64_t fib(int64_t n) {  
    if (n < 2) {  
        return n;  
    } else {  
        int64_t x, y;  
        #pragma omp task shared(x,n)  
        x = fib(n-1);  
        #pragma omp task shared(y,n)  
        y = fib(n-2);  
        #pragma omp taskwait  
        return (x + y);  
    }  
}
```

Wait for the two tasks
to complete before
continuing.

Fibonacci in OpenMP

- OpenMP provides many **pragma directives** to express common patterns
 - ▣ e.g. **parallel for** for loop parallelism
 - ▣ e.g. **reduction** for data aggregation
 - ▣ e.g. directives for scheduling and data sharing
- OpenMP supplies a variety of **synchronization constructs**,
 - ▣ e.g. barriers, atomic updates, mutual-exclusion (mutex) locks

CONCURRENCY PLATFORMS

- Pthreads (& WinAPI Threads)
- Threading Building Blocks
- OpenMP
- **Cilk**



Cilk

- Award-winning multithreading language developed at MIT
- Linguistic **extensions to C/C++** to support **fork-join** parallelism
- Features a provably efficient **work-stealing** scheduler
- Provides a **reducer** linguistic interface for parallelizing code with global variables

OpenCilk

- This course will be using the **OpenCilk** compiler
 - ▣ OpenCilk is based on **LLVM**
 - ◆ developed at MIT by Tao B. Schardl, William S. Moses, and Charles E. Leiserson
 - ▣ OpenCilk's **compiler** can generally produce better code than can other compilers for parallel-language constructs
 - ▣ The OpenCilk **runtime** system is based on Cheetah
 - ◆ developed by I-Ting Angelina Lee at Washington University in St. Louis
 - ▣ Tools: the **Cilksan** race detector and the **Cilkscale** scalability analyzer

Nested Parallelism in Cilk

```
int64_t fib(int64_t n) {  
    if (n < 2)  
        return n;  
    int64_t x, y;  
    cilk_scope {  
        x = cilk_spawn fib(n-1);  
        y = fib(n-2);  
    }  
    return (x + y);  
}
```

The named **child** function may execute in parallel with the **parent** caller.

Control cannot pass this point until all spawned children have returned.

Loop Parallelism in Cilk

Example: in-place
matrix transpose

$$\begin{matrix} \begin{pmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nn} \end{pmatrix} & \xrightarrow{\hspace{1cm}} & \begin{pmatrix} a_{11} & a_{21} & \dots & a_{n1} \\ a_{12} & a_{22} & \dots & a_{n2} \\ \vdots & \vdots & \ddots & \vdots \\ a_{1n} & a_{2n} & \dots & a_{nn} \end{pmatrix} \\ A & & A^T \end{matrix}$$

The iterations of a
cilk_for loop
execute in parallel

```
// indices run from 0, not 1
cilk_for (int i=1; i<n; ++i) {
    for (int j=0; j<i; ++j) {
        double temp = A[i][j];
        A[i][j] = A[j][i];
        A[j][i] = temp;
    }
}
```

Serial Semantics

Cilk source

```
int64_t fib(int64_t n) {  
    if (n < 2)  
        return n;  
    int64_t x, y;  
    cilk_scope {  
        x = cilk_spawn fib(n-1);  
        y = fib(n-2);  
    }  
    return (x + y);  
}
```



serial projection

```
int64_t fib(int64_t n) {  
    if (n < 2) {  
        return n;  
    } else {  
        int64_t x, y;  
        x = fib(n-1);  
        y = fib(n-2);  
        return (x + y);  
    }  
}
```

The **serial projection** of a Cilk program is always a legal interpretation of the program's semantics.

Remember, Cilk keywords **grant permission** for parallel execution. They do not **command** parallel execution.


To obtain the serial projection:

```
#define cilk_for for  
#define cilk_spawn  
#define cilk_scope
```


Scheduling

- Cilk allows the programmer to express **logical parallelism** in an application
- The Cilk **scheduler** maps the executing program onto the processor cores dynamically at runtime
- Cilk's **work-stealing** scheduling algorithm is provably efficient


```
int64_t fib(int64_t n) {  
    if (n < 2)  
        return n;  
    int64_t x, y;  
    cilk_scope {  
        x = cilk_spawn fib(n-1);  
        y = fib(n-2);  
    }  
    return (x + y);  
}
```




Cilk Platform




Serial Testing




Alternative Serial Testing



Parallel Testing



Scalability Analysis



Summary

- Processors today have **multiple cores**, and obtaining high performance requires **parallel programming**
- Programming directly on processor cores is **painful** and **error-prone**
 - Concurrency platforms simplify programming for you
- **Cilk** simplifies parallel programming, and performs **provably** efficiently
- Project 2: Parallel simulation & rendering using Cilk